



MSc in Computer Science 2017-18

Project Dissertation

Project Dissertation title: Reflectance Transformation Imaging

Term and year of submission: Trinity Term 2018

Candidate Name: Johannes Bernhard Goslar

Title of Degree the dissertation is being submitted under: MSc in Computer Science

Abstract

This project presents a new web-based Reflectance Transformation Imaging (RTI) viewer. This project will also explore that software's potential in practical applications, such as in the study of historical objects. It is based on a novel plugin structure, allowing easy extension by third parties. Additionally included is a new file format (specification) for RTI images, enabling better compression and web compatibility; and novel viewer functionality like overlays. Fundamentally this project was inspired by a desire to improve the experience of users and developers of RTI viewer(s).

Acknowledgements

I want to thank following people for being part of the journey that culminated in this project. Roland for starting everything by forcing me to work through that Visual Basic 2005 book. Olivier and Søren for their guidance and pushing me into international territories. Jassim and Stefano for rekindling my interest in Computer graphics. Taylor for providing real users. Christiane and Elsie for always being motivational and supportive.

Contents

1	Introduction	1
2	Related Work	1
2.1	RTI Theory and Workflows	2
2.2	Fileformats	3
2.3	RTI Viewers	4
3	Requirements	4
4	Design	6
4.1	Fileformat	6
4.2	Architecture	7
4.2.1	State Driven Plugins	8
4.3	Renderer Stack	8
5	Implementation	10
5.1	Overview	10
5.2	Libraries	10
5.3	Base	19
5.4	Hooks	20
5.5	BTF File	21
5.6	State Management	22
5.7	Renderer Stack	23
5.8	Texture Loader	26
5.9	Plugins	27
5.9.1	Base Plugin	30
5.9.2	BaseTheme, RedTheme and BlueTheme Plugin	30
5.9.3	TabView Plugin	31
5.9.4	SingleView Plugin	31
5.9.5	Converter Plugin	31
5.9.6	PTMConverter Plugin	34
5.9.7	Renderer Plugin	35
5.9.8	PTMRenderer Plugin	36
5.9.9	LightControl Plugin	38
5.9.10	Rotation Plugin	38
5.9.11	Zoom Plugin	40
5.9.12	QuickPan Plugin	40
5.9.13	Paint Plugin	41
5.9.14	Bookmarks Plugin	41
5.9.15	Notes Plugin	41
5.9.16	ImpExp Plugin	45
5.9.17	Settings Plugin	45
5.10	Targets	46
5.10.1	Electron	46
5.10.2	Web	46

5.10.3	Hosted	48
5.10.4	Embeddable	48
5.10.5	Mobile Phones	49
6	Results	49
6.0.6	Performance	49
6.1	Accuracy	52
6.2	Rollouts and Testing	52
7	Discussion	54
7.1	Novelties	54
7.2	Future Work	54
7.3	Community Outreach and Onboarding	56
8	Conclusion	56
A	BTF File Format	57
A.1	File Structure	57
A.2	Manifest	58
A.3	Textures	58
A.4	Options	59
A.5	Layers	60
A.6	State	60

List of Figures

1	RTI Workflow	2
2	RTI in Action	2
3	RTI Overview	3
4	Unpacked BTF file	8
5	Class and Component Overview	9
6	MobX Flow	14
7	WebGL compatibility	15
8	gl-react-cookbook example	16
9	Zoom Component	19
10	MobX actions	23
11	MobX state tree	24
12	View Comparison	32
13	Converter UI	33
14	Normal Rendering	39
15	LightControl Plugin	39
16	Zoom Plugin	40
17	Paint Plugin	42
18	Bookmark Plugin	43
19	Notes Plugin	44
20	Settings Plugin	45
21	Electron App	47

22	Embedded Version	48
23	RTI Mobile	50
24	Load time comparison	51
25	Warrior Comparison	53
26	Tablet Comparison	53

List of Tables

1	Performance Comparison	49
2	RMSE for Accuracy Comparison	52

1 Introduction

Reflectance Transformation Imaging (RTI) was invented by Tom Malzbender, Dan Gelb and Hans Wolters, research scientists at Hewlett-Packard Labs[28]. It was originally termed Polynomial Texture Mapping (PTM). It is a method of computational photography with great potential for classical archaeology. RTI images (RTIs) are created from multiple digital photographs of an object. A fixed camera position is used in conjunction with a movable light source, or multiple immovable light sources, hitting the object from different positions. Different shadows and highlights result from each light position.

An RTI processing software takes these images and calculates the object's surface per pixel, essentially creating a 2D photograph with embedded reflectance information.

This file can then be viewed with the help of an RTI viewing software, allowing the object to be studied remotely via the user's computer, instead of needing physical access to the object. The software is also able to reveal enhanced or previously unobservable details, e.g. colour, shape, markings or depth of carved lines, which the naked eye could not pick up.

Ancient historians studying material objects benefit from RTI because this software shows how objects were created and subsequently changed. For example captured RTIs of the wooden remains of Roman wax tablets can reveal how they have been inscribed and reinscribed. For statues, particularly Roman imperial portraits RTI can provide evidence for deliberate re-carving of a condemned emperor into his successor.

For video examples and diagrams, the Cultural Heritage Imaging Institute provides a great overview[12].

Other applications of RTI are video games, where RTI can provide self-shadowing for objects, which normal maps could not, though more modern techniques largely replaced RTI in this context.

They are becoming increasingly relevant in a physical based rendering (PBR) context, where RTIs can provide better physical information for objects.

This dissertation is primarily focusing on the application of RTI for ancient historians, but makes provisions for relevance in a PBR context in the future.

oxrti in this document is referring to this project's implementation of an RTI viewer.

2 Related Work

This section explores the related work inside the RTI area and explains the underlying principles. The project is currently only concerned with the viewing part, so



Figure 1: RTI workflow, courtesy of CHI[13].

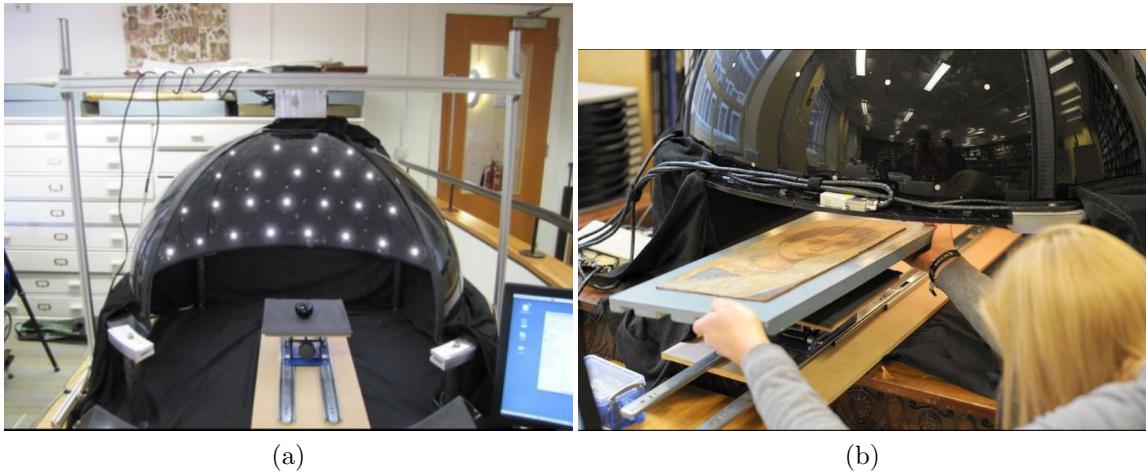


Figure 2: (a) shows the spread out light sources inside an RTI dome. (b) shows an object being positioned. Images from Piquette and Crowther.[34]

that will be covered in the most depth.

2.1 RTI Theory and Workflows

The RTI workflow contains three steps (see Figure 1) in our context (with the hunt for the physical object excluded). The object is placed centrally under some capturing device (like Figure 2). A camera is statically mounted on top. The lights are set up so that all captured images will have a different light position. From these images the RTI coefficients are calculated (see Figure 3).

For example the PTM LRGB format[28] stores 9 coefficients per texel: $a_0 - a_5$ for calculating the varying luminance and R, G, B for colours, the resulting rendering is calculated in the following way per pixel:

$$\begin{aligned} R(u, v) &= L(u, v)R_n(u, v); \\ G(u, v) &= L(u, v)G_n(u, v); \\ B(u, v) &= L(u, v)B_n(u, v); \end{aligned}$$

With u and v being the texture coordinates. R_n , G_n and B_n being the raw unscaled

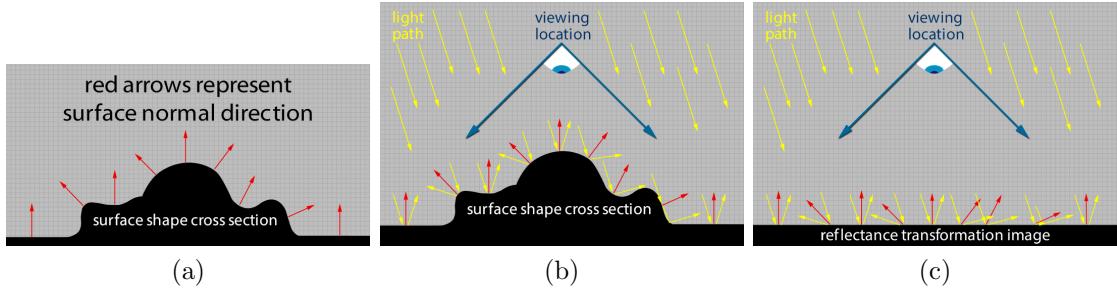


Figure 3: (a) shows the to be captured object’s normals. (b) shows the object with a single viewpoint and light source (in the capturing dome the eye would be camera). (c) visualizes the RTI viewing process. Images from CHI[12].

input colours and $L(u, v)$ calculated as following:

$$L(u, v, l_u, l_v) = a_0(u, v)l_u^2 + a_1(u, v)l_v^2 + a_2(u, v)l_u l_v + a_3(u, v)l_u + a_4(u, v)l_v + a_5(u, v)$$

Where a_0-a_5 are the previously fit coefficients and l_u and l_v are x and y components of the light vector projected into the texture plane. The coefficients are calculated as part of the RTIBuilder process, for details refer to [28]. Other file formats will have slightly different sets of coefficients, but the same principles hold, e.g. PTM RGB has 6 luminance coefficients per colour channel and no fixed colours.

2.2 Fileformats

The most comprehensive overview on the current state of the art is done by the American Library of Congress (LoC) as part of its digital preservation effort, with sections on the PTM[8] and RTI[9] formats. The current PTM specification by Malzbender et al[28] from 2001 specifies 9 different formats, of which “the most commonly used encoding formats are [...] PTM_FORMAT_LRGB (aka LRGB PTM) and PTM_FORMAT_RGB (aka RGB PTM)”[8], so most effort was spent to support these inside the oxrti implementation. Both formats begin with a metadata block consisting of 4 elements separated by newlines:

- The version prefix, which should be ‘PTM_1.2’ for recent files.
- The format string, e.g. PTM_FORMAT_LRGB.
- Image size, the dimensions in pixels.
- Scales and biases, to scale and bias the coefficients before calculation.

This metadata block is followed by the raw texel data in forms depending on the format All in reversed scanline order. PTM_FORMAT_RGB stores a_0-a_5 per texel in separated RGB blocks. PTM_FORMAT_LRGB stores a_0-a_5 per texel first, then followed by RGB blocks of a single coefficient per texel. The destructure of these is presented in detail in section 5.9.6. Different embedded JPEG compression styles are specified as well, but due to their relatively small popularity and increased complexity, the initial oxrti version has omitted them.

The RTI fileformat specification by Corsini et al.[10] is a newer format supporting Hemi-Spherical Harmonics (HSH) in addition to PTM’s biquadratic polynomials, as HSH “handles shiny surfaces and specular highlights (bright spots of light appearing on shiny objects) better than PTM, which yields matte images.”[9]. One additional difference is the support of metadata as an XMP packet at the end of the file. As most locally available data is in PTM format, support for the RTI format is initially omitted for this project, but it will be easy to add support in the future.

2.3 RTI Viewers

The following RTI Viewers are currently available:

PTM Viewer from HP. The initial viewer application, available at [22]. Downloads for Windows, MacOS, HP-UX and Linux. Only rudimentary functionality of changing the light direction. The MacOS version did not run on the author’s computer. No source code available.

RTIViewer available from the CHI at [13]. Most widely used viewer. Downloads for Windows and MacOS. Source code only available per request: “This software is available under the Gnu General Public License version 3. If you wish to receive a copy of the source code, please send email to info@c-h-i.org.” PTM and RTI formats supported. Features: Light control, multiple rendering modes, bookmarks.

InscriptiFact Viewer by the University of Southern California. Offered in two flavours, a downloadable Java program[24] and a embedded Java applet[25] (not working as of August 28, 2018). No source code available. Features: Light control, multiple rendering modes.

webRTIViewer by the Visual Computing Laboratory. Download at [38]. Support for a proprietary format, requiring pre-processing with a C++ command line utility. Source code included in the download. Features: Light control.

OxRTI SVG Viewer by Christopher Ramsey. SVG based web viewer available at [32], developed at the same time as this project. Requiring pre-processed PTMs. No source code available at the moment, though a future publication is announced. Features: Light control, multiple rendering modes, comments and painting.

3 Requirements

In preparation for the project the author had multiple informal conversations with consenting RTI users from different Oxford departments in relation to their current usage of RTI viewing software and wishes for a modernized software. The prevailing wish was for a unification of the software that is used, most researchers had to export screenshots from the RTI software, to then add annotations and drawings

in their graphical editing software. But in their graphical editing software all the advantages of an RTI image were then lost. Following these discussions, the following requirements were decided upon. The assumption was made that this project be considered an exploratory piece of work to show possible new directions for RTI softwares.

The first overall goal is the main focus on web technologies. The new viewer should be available on as many platforms as possible and allow for easy distribution and usage, as well as a comfortable developer experience. Because of this, the next priority was a web-supported fileformat. No browser currently supports RTI or PTM files natively, so some conversion needs to take place. The requirement of an additional conversion software would be diametrical to the ‘everything web based’ goal though, so the browsers need to be able to also handle the conversion and then support some kind of export. The easiest way to do so was to develop a newer file format with following features:

1. Support for embedding/consuming the PTM[8] fileformat.
2. Future support for embedding/consuming the RTI[9] fileformat.
3. Extended metadata support.
4. Support for high resolutions.
5. Support for higher bitdepths per pixel than the 8 of PTM/RTI for future HDR applications.
6. Easy exchange between multiple researchers and computers.

For the viewer the following features were deemed essential:

7. Compatible with all major operating systems and/or web browsers.
8. Light Controls to change the position of the virtual light source.
9. Multiple rendering modes apart from plain LRGB or RGB.
10. Quick navigation functionality.
11. Annotations to allow for further textual information relating to a part of the viewer object.
12. Paintable layers to further clarify surface details.

In addition these non-functional requirements were decided upon:

13. Free Software, the implementation should be available for everyone to change and distribute. No email-walling, instead embracing jump-in interactions between multiple parties.
14. Suitable for new developers: for scientists for research purposes, or students for educational purposes.
15. Plugin support to allow independent additions or modifications to the core software.

16. Good developer experience.
17. Adequate performance, at least keeping up with current implementations.
18. Easy installation for researchers.
19. Fast responses to user interactions.
20. Reasonable file sizes for instant transfer/viewing.
21. Preservable software and BTF files.

4 Design

4.1 Fileformat

Although the LoC deems the PTM format “relatively transparent. It can be viewed in plain text editors. [...] A very simple program would present these numbers in a human-readable form.”[8] the author thinks that another implementation could be even more transparent and easy for people to understand. One simple initial step is to move to multiple files for representing the coefficients, each being a valid image (BMP or PNG) themselves. This enables easy examination, simplifies texture loading in a web context (as PNG textures are supported) and allows easy interactions with other tools. Having multiple files anyway, the textual data description and metadata can also be stored in an extra plain text file, making it editable with standard tools. How to easily distribute multiple files though - just archive (zip) them and distribute as one. Web browsers do not unpack zips inside their JavaScript run times natively, but libraries[26] exist for that functionality. All this multi file handling is abstracted away from the user, they will just handle a single `.btf.zip` file, but provides access if interested. Having the ability to add extra files to the archive also allows for easy extensions, for example the oxrti implementation adds the layer and state saving functionalities, which are stored as extra entries. For an unpacked example see Figure 4. Having the web as first target, JSON files are the natural match to describe the contained data. The data specification’s main goal was future extensibility and coverage of all existing formats, the developed specification is available standalone at the end of this document (Appendix A) to accommodate easy redistribution and improve accessibility. One fairly self-explanatory example of converter output:

```

1  {
2      "name": "tablet2_LRGB",
3      "data": {
4          "width": 512,
5          "height": 512,
6          // the channel model declares how the channels are combined to form
7          // the
8          // resulting picture
9          "channelModel": "LRGB",

```

```

9   "channels": {
10     // each channel declares how it should be calculated, RTIpoly6 is
11     // the calculated described in the related work section
12     "L": {
13       "coefficientModel": "RTIpoly6",
14       "coefficients": {
15         "a0": { "format": "PNG8" },
16         "a1": { "format": "PNG8" },
17         "a2": { "format": "PNG8" },
18         "a3": { "format": "PNG8" },
19         "a4": { "format": "PNG8" },
20         "a5": { "format": "PNG8" }
21       }
22     },
23     "R": {
24       // flat is just taking the straight value from the texture
25       "coefficientModel": "flat",
26       "coefficients": {
27         "a0": { "format": "PNG8" }
28       }
29     },
30     "G": {
31       "coefficientModel": "flat",
32       "coefficients": {
33         "a0": { "format": "PNG8" }
34       }
35     },
36     "B": {
37       "coefficientModel": "flat",
38       "coefficients": {
39         "a0": { "format": "PNG8" }
40       }
41     },
42     // format extra is channelModel dependant and of free JSON form
43     "formatExtra": {
44       "scales": [2, 2, 2, 2, 2, 2],
45       "biases": [149, 151, 68, 77, 77, 0]
46     }
47   }
48 }
```

4.2 Architecture

The primary goal for the implementation was to provide many small parts, which all should be replaceable and flexible to allow the easy addition of any new wished-for features and the integration of new rendering models. This requires abstractions in two areas: a flexible plugin system supporting the transparent integration of new



Figure 4: Unpacked BTF file, manifest.json is the standardised data description, oxrti_* are extensions. The L folder contains textures for the single coefficients in grey scale format.

user interfaces and a flexible rendering mechanism, as the different rendering models should not each be required to implement the viewing features like panning.

4.2.1 State Driven Plugins

The first can be solved as following: Single code units ('Plugins') containing all their functionality and only integrating with the rest of the program through an indirect system ('Hooks'), where plugins can declare their additions to a specific hook - be it a notification hook, where plugins register a function to be called when an event occurs; or a configuration hook, with which plugins extend some core functionality with their additions (e.g. adding a user interface component at some place). The problem of reactive data exchanges between plugins (e.g. light position change for the rendering algorithms) and data bindings (for the user interfaces) remains to be solved however. 'Hooking' every variable read and write would result in excessive code bloat (however, *that would still be a more welcome solution than just writing variables). Instead an observer system should be used. Variable changes should be tracked transparently and reading functions/components be notified of changes automatically.

This proposes the introduction of a state-driven plugin conforming to the Model-View-Controller pattern. Each plugin should explicitly declare what its state is (the model), have some specific code changing that state (the controller) and have one or more user facing components (the view(s)). With this the observations are one directional because they only go from state to models, which is far easier to track.

For the current oxrti implementation, all Plugins and components are depicted in Figure 5 and the concrete implementation is shown in section 5.9.

4.3 Renderer Stack

A more traditional frame based rendering stack would likely either re-render everything every frame (like a realtime game) or fully render all shaders on any input change. But in conjunction with the observation system above it is possible to implement a more efficient one. Rendering plugins register themselves at the appropriate

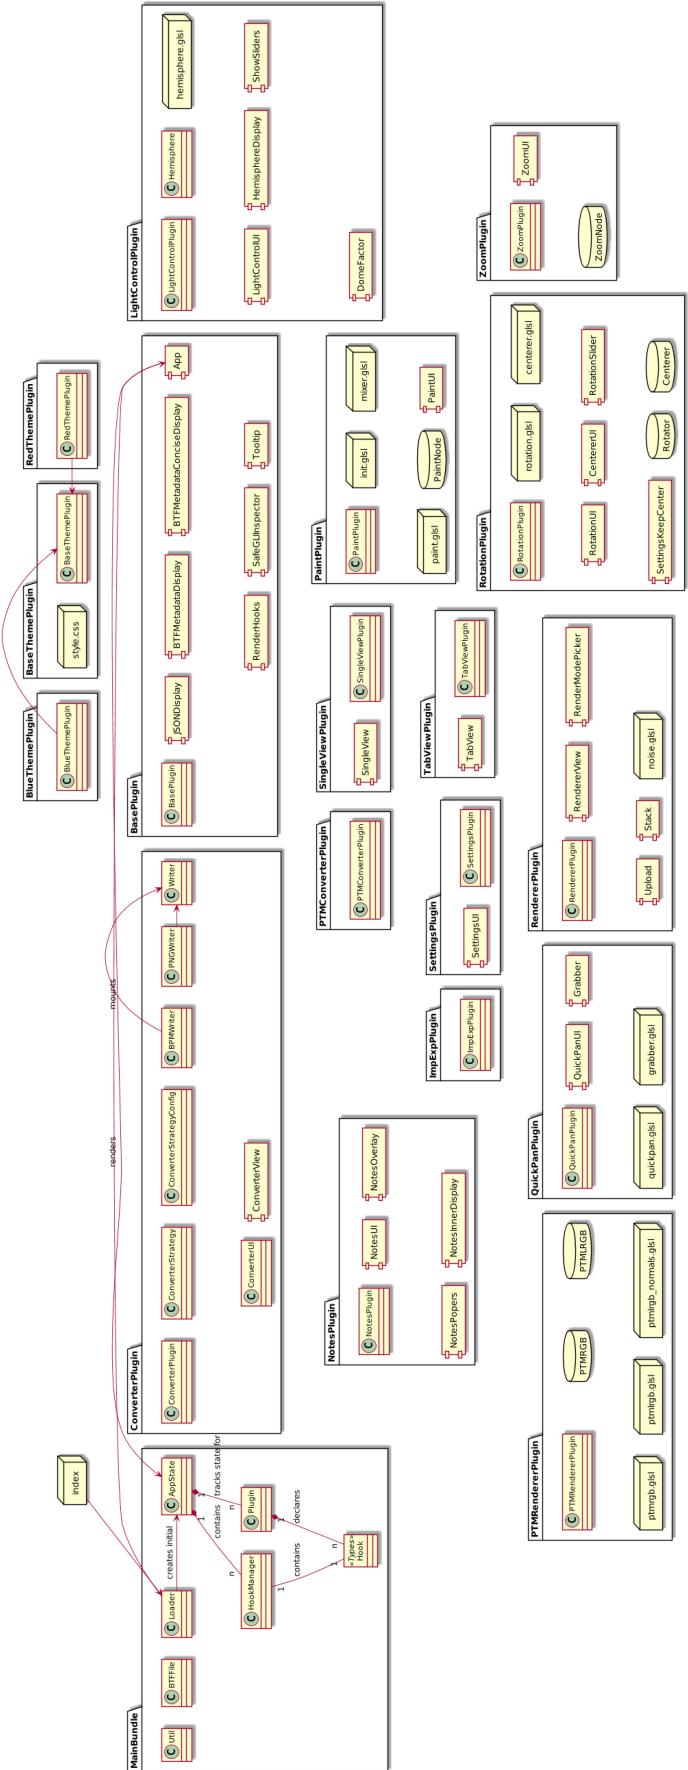


Figure 5: All classes and components of the implementation. Inheritance arrows from each Plugin to the Plugin base class have been omitted for sake of clarity. Boxes are files, usually `html/css/gsl`. Cylinders are rendering stack nodes. Rectangles with two connectors on the left are user interface components.

hook with a specific priority (highest rendered first). When a rendering output changes, the successor node inside the rendering stack is notified of the change and performs a re-rendering on its own. With this, expansive operations like the model rendering can be cached on lower levels of the stack. For example the zoom plugin merely transforms these cached outputs. The concrete implementation is discussed in section 5.7.

5 Implementation

5.1 Overview

Since this section explains the current implementation of the developed tool set, it not only aims to fulfill the dissertation's requirements, but also to help users who want to understand the underlying systems and prepare them for potentially joining the development effort. Although abridged code extracts are of their current state for thesis submission, and the main principles will remain, nevertheless future readers are welcome to consult the actual source code if any discrepancies arise or reexport the document.

Firstly the main libraries are briefly explained where they are relevant to the program. Secondly, the mostly abstract plugin architecture is shown. Thirdly the main plugins are presented and last the delivery processes to the end users are described.

All implementation files are contained and delivered inside a single git repository, which is freely available online: <https://github.com/ksjogo/oxrti>. All following file paths are relative to that repository's root. All future development will be immediately available there and the current compiled software version is always fed automatically from it into the hosted version at <https://oxrtimaster.azurewebsites.net/api/azurestatic>.

5.2 Libraries

TypeScript

The official header line of TypeScript gives reasons why it was picked for this project: “TypeScript is a typed superset of JavaScript that compiles to plain JavaScript. Any browser. Any host. Any OS. Open source.”[44] Which fits requirements 13, 7. Whereas plain JavaScript would have allowed slightly easier initial on-boarding and maybe easier immediate code ‘hacks’, TypeScript will provide better stability in the long term and a quite improved developer experience (requirement 16) in the long run. With the full typed hook system (compare section 5.4) it ensures that a compiled plugin will not have runtime type problems, reducing the amount of switching between code editor and the running software. The whole project is set up

in a way to fully embrace editor tooling. Visual Studio Code[45] and Emacs[20] are the ‘officially’ tested editors of the project. Code is recommended as it will support all developer features out of the box. The installation of the tslint[42] plugin[43] is recommended to keep a consistent code style, which is configured within the *tslint.json* file. Most importantly TypeScript adds type declarations (and inference) to JavaScript, e.g.:

```
1 const thing = function (times: number, other: (index: number) => boolean)
  ↵ { ... }
```

would define *thing* as a function, taking a numbers as first argument and another function (taking a number as first parameter and returning a boolean) as second argument. The other most used TypeScript features inside the codebase are Classes[6], Decorators[14] and Generics[17], which will be discussed at their first appearance inside the code samples.

React

The two main points on React’s official website are “Declarative” and “Component Based” [36], which is best shown by an extended example from their website, which exemplifies multiple patterns found through the oxrti implementation. The most important concept is the jump from having a stateful HTML document, which the JavaScript code is manipulating directly, e.g.:

```
1 document.getElementById('gsr').innerHTML=<p>You shouldn't do this</p>"
```

Which is diametric to requirements 14 and 16 as it would require developers to manually keep track of all data cross-references (e.g. the pan values having to automatically adapt to the current zoom level). A declarative approach instead allows much better and easier implemented reactivity and better performance (requirements 17 and 19) as the necessary changes can be tracked and components can be updated selectively.

```
1 // a class represents a single component
2 class Timer extends React.Component {
3   // the parent component can pass on props to it
4   constructor(props) {
5     super(props);
6     this.state = { seconds: 0 };
7   }
8
9   tick() {
10    // the state is updated and the component is automatically
11    ↵ re-rendered
11   this.setState(prevState => ({
12     seconds: prevState.seconds + 1
13   }));
14 }
```

```

16 // called after the component was created/added to the browser window
17 componentDidMount() {
18     this.interval = setInterval(() => this.tick(), 1000);
19 }
20
21 // called before the component will be deleted/removed from the browser
22 // window
22 componentWillUnmount() {
23     clearInterval(this.interval);
24 }
25
26 // the actual rendering code
27 // html can be directly embedded into react components
28 // {} blocks will be evaluated when the render method is called
29 // which will happen any time the props or its internal states updates
30 render() {
31     return (
32         <div>
33             Seconds: {this.state.seconds}
34         </div>
35     );
36 }
37 }
38
39 // mountNode is a reference to a DOM Node
40 // the component will be mounted inside that node
41 ReactDOM.render(<Timer />, mountNode);

```

In conjunction with mobx and TypeScript no classes are used for React components though, but instead Stateless Functional Components ('SFCs'[23]). These SFCs are plain functions, only depending on their passed properties:

```

1 function SomeComponent(props: any) {
2     return <p>{props.first} {props.first}</p>
3 }

```

This component could then be used by:

```

1 <SomeComponent first="Hello" second="World"/>

```

This component systems allows the plugins to define some components and then 'link' them into the program via the hook system, which will be explored later.

MobX

Its main tagline is "Simple, scalable state management"[30]. An introductory overview is shown in Figure 6. Broadly speaking MobX introduces observable objects. Instead of the aforementioned DOM handling or property passing inside React trees, components can just retrieve their values from the observable objects and will be automatically refreshed if the read values change. This for example makes the

implementation of the QuickPan plugin extremely easy, as it can just read the zoom, pan, etc. values of the other plugins and will automatically receive all updates without any further manual observation handling.

mobx-state-tree

“Central in MST (mobx-state-tree) is the concept of a living tree. The tree consists of mutable, but strictly protected objects”[31] This allows the implementation to have one shared state tree which can be used to safely access all data. All nodes inside the state tree are MobX observables. A simple tree with plain MST would look like this:

```
1 // define a model type
2 const Todo = types
3   .model("Todo", {
4     // state of every model
5     title: types.string,
6     done: false
7   })
8   .actions(self => ({
9     //methods bounds to the current model instance
10    toggle() {
11      self.done = !self.done
12    }
13  }))
14 // create a tree root, with a property todos
15 const Store = types.model("Store", {
16   todos: types.array(Todo)
17 })
```

This syntax was deemed to be convoluted, as it is a lot more complex than standard JavaScript/TypeScript classes, which were introduced by the ES6 version, as shown in the React description above and thus being in conflict with requirement 14.

classy-mst

There is an option to use a more traditional syntax instead though, with the classy-mst library, with which the example above becomes[7]:

```
1 const TodoData = types.model({
2   title: types.string,
3   done: false
4 });
5
7 class TodoCode extends shim(TodoData) {
8   @action
9   toggle() {
```

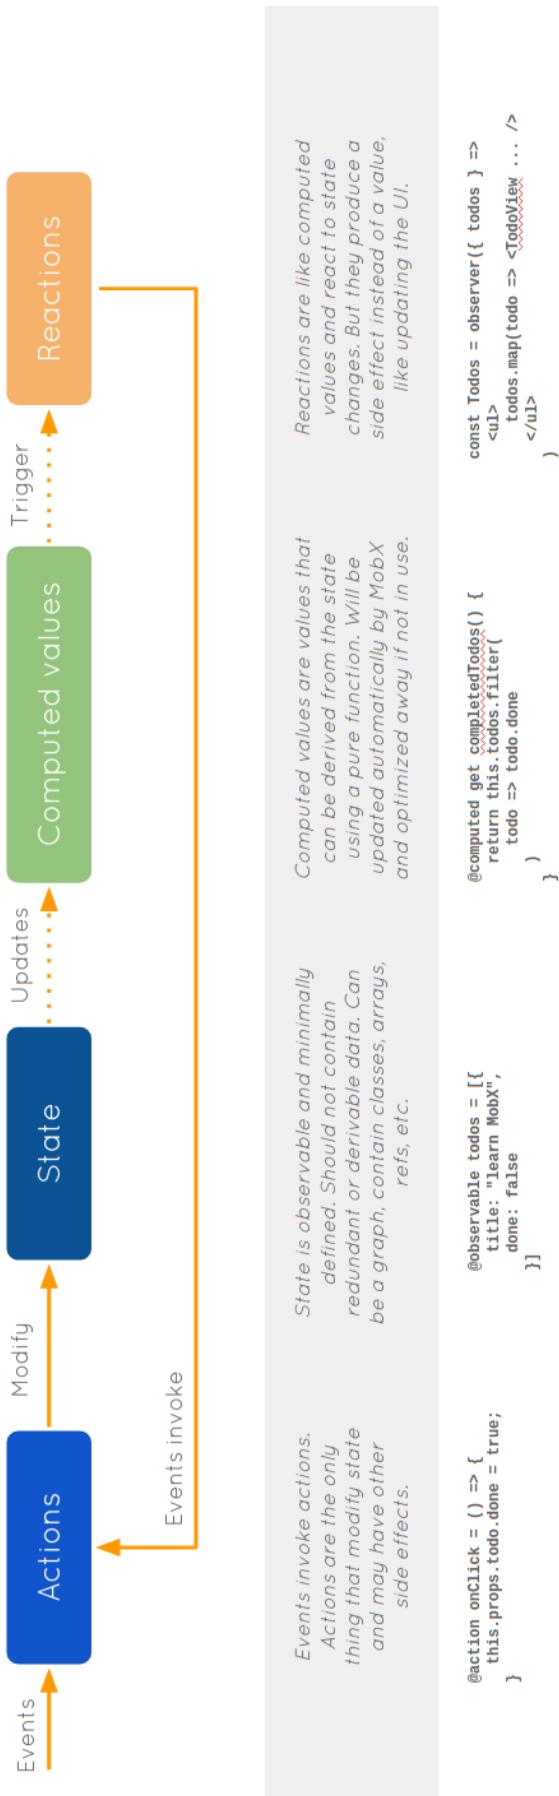


Figure 6: Taken from Weststrate[30]. Actions in the oxrti context are most often initially user actions, which are then calling into plugins to change the state. The state is mostly encapsulated on a plugin basis with usage of the mobx-state-tree library, which also encapsulates most computed values. Reactions are most often the previously discussed React components.

Desktop		Mobile					
Feature		Chrome	Edge	Firefox (Gecko)	Internet Explorer	Opera	Safari
Basic support		9	(Yes)	4.0 (2.0)	11	12	5.1
WebGL 2		56	No support	51 (51)	No support	43	No support

Desktop		Mobile					
Feature		Chrome for Android	Edge	Firefox Mobile (Gecko)	IE Mobile	Opera Mobile	Safari Mobile
Basic support		25	(Yes)	4	No support	12	8.1
WebGL 2		?	?	?	?	?	?

Figure 7: WebGL compatibility as from the Mozilla Developer network[41].

```

10      this.done = !this.done;
11  }
12 }
13
14 const Todo = mst(TodoCode, TodoData, 'Todo');

```

Weststrate, the original author of MobX initially was sceptic of this syntax[4] as it was changing the semantics of ES6 classes, as classy-mst’s methods will be automatically bound to the instance. This boundness is an advantage for this implementation though, as the hook configurations can just refer to `this.someMethod` instead of `this.someMethod.bind(this)`. The `@action` is a decorator, enabling the following method to change the state/properties of the model, as MST prohibits that by default. Reactions/View updates will only happen after the outermost action finished executing.

WebGL

The increasing support of the WebGL stack is the main reason why it is now feasible to implement a full RTI software stack with plain web technologies, as it “enables web content to use an API based on OpenGL ES 2.0 to perform 3D rendering in an HTML `<canvas>` in browsers that support it without the use of plug-ins.”[49] OpenGL ES 2.0 likeness means that (most importantly) shaders are supported, allowing the implementation to be split up into multiple shaders with single responsibilities. For details refer to section 5.7. While preserving compatibility and requirement 7 WebGL 2 support is sadly not widespread enough to fully rely on yet (compare Figure 7), as it is currently estimated at 50% of all devices[47]. Potential improvements when WebGL 2 is more widely supported or in conditional plugins are discussed in section 7.2. One notable limitation of WebGL is `MAX_TEXTURE_IMAGE_UNITS`, the maximum amount of bound textures inside a single shader, which in most implementations is 16[48], whereas the standard OpenGL implementations are likely to have a limit of 32. This is influencing the BTF file format, as for example in the PTM RGB use case a total of 18 coefficients exist, which now need to be bundled up somehow into maximum 16 textures, if the calculations should be done inside a single shader. It is also limiting the amount of layers of the Paint plugin, as these also consist of bound textures. Apart from the shaders,

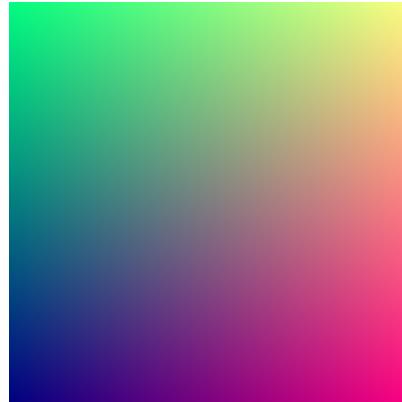


Figure 8: RGBA texture, with R and G according to their respective u or v texture coordinate. From [18].

which are written in the OpenGL ES Shading Language[46] and the the texture loader (section 5.8), no direct WebGL code is necessary nor used anywhere inside the implementation, as the gl-react library is abstracting it neatly for use from the MobX/React environment.

gl-react

“Implement complex effects by composing React components.”[39] is the main use of the gl-react library. A minimal component, adapted from the gl-react-cookbook looks like[39]:

```

1 const shaders = Shaders.create({
2   helloGL: {
3     frag: GLSL`  

4       precision highp float;  

5       varying vec2 uv;  

6       void main() {  

7         gl_FragColor = vec4(uv.x, uv.y, 0.5, 1.0);  

8       }`  

9     }
10   });
11
12 export default class Example extends Component {
13   render() {
14     return (
15       <Surface width={300} height={300}>
16         <Node shader={shaders.helloGL} />
17       </Surface>
18     );
19   }
20 }
```

Which would result in a display like Figure 8. gl-react’s is not a 3D engine, so no

objects are to be created or scene graph managed, instead the oxrti implementation can concentrate on solely providing the necessary shaders. gl-react's default node size is taken from the parent surface size. The surface size will be dependent on the user running the program and his browser windows, which makes it undesirable as details would be lost, if the BTF provided more detail, so the processing Node sizes are usually set to the BTF resolution or higher.

Webpack

Webpack is used to bundle the implementation into single files as it is “a bundler for javascript and friends. Packs many modules into a few bundled assets.”[50] A more detailed discussion on the targets is in section 5.10. Broadly speaking Webpack loads the source code inside the *src* directory according to the loaders defined inside the *webpack.config.js* file, analyses their dependencies and then bundles them together. This makes it possible to have a dependency tree spanning 26184 packages from npm, but still providing a single bundled application file only 1.5 megabyte large (data as of August 28, 2018). It also allows the dynamic plugin structure by bundling the plugins into a dynamic ‘context’ from which single plugins can be loaded at runtime.

Electron

Electron is used to “build cross-platform desktop apps with JavaScript, HTML, and CSS”[15] While theoretically not necessary to fulfill most requirements, as the implementation is compatible with all modern web browsers, an additional standalone executable provides some advantages:

- It is possible to add a more traditional menu-based interface, which the browser version could not support.
- Stable development environment, as electron-devtools-installer is used to provide relevant extensions (React devtools, MobX devtools) by default and the hot reloading is reliably tested, which together form a good developer experience (requirement 16)
- It allows to preserve the software in a usable, contained state, not relying on the user also having a compatible web browser in the future.
- It allows future development to more directly access resources of the host machine, e.g. the normal OpenGL stack could be used for calculating the coefficients, as it is less resource constrained compared to the WebGL stack.

MaterialUI

MaterialUI is succinctly described by “React components that implement Google’s Material Design.”[29]. MaterialUI’s component are used throughout the app for

styling the components, making the use of custom CSS largely unnecessary apart from minor positioning fixes. For example the Zoom component is defined as:

```

1 // Card, CardContent, Tooltip and Button are all components provided by
2   ↳ MaterialUI.
3 // this refers to the Zoom Plugin's controller which
4 // the content will be automatically refreshed if the referred values
5   ↳ change
6 const Zoom = Component(function ZoomSlider (props) {
7   return <Card style={{ width: '100%' }} >
8     <CardContent>
9       <Tooltip title='Reset'>
10         <Button onClick={this.resetZoom} style={{ marginLeft:
11           ↳ '-8px' }}>Zoom</Button>
12       </Tooltip>
13       <Tooltip title={this.scale}>
14         <Slider value={this.scale} onChange={this.onSlider}
15           ↳ min={0.01} max={30} />
16       </Tooltip>
17       <Tooltip title='Reset'>
18         <Button onClick={this.resetPan} style={{ marginLeft:
19           ↳ '-11px' }}>Pan</Button>
20       </Tooltip>
21       <Tooltip title={this.panX}>
22         <Slider value={this.panX} onChange={this.onSliderX}
23           ↳ min={-1 * this.scale} max={1 * this.scale} />
24       </Tooltip>
25       <Tooltip title={this.panY}>
26         <Slider value={this.panY} onChange={this.onSliderY}
27           ↳ min={-1 * this.scale} max={1 * this.scale} />
28       </Tooltip>
29     </CardContent>
30   </Card>
31 })

```

It would result in a display like Figure 9.

Misc

Further libraries of note are:

electron-webpack[16] is providing the bridging between Webpack and Electron, its config is expanded by the *webpack.renderer.additions.js* and *webpack.renderer.shared.js* files.

pngjs[40] is providing in-browser bitwise png manipulation, required in the converter.

jszip[26] is providing in-browser zip file manipulations, which are fundamental for the BTF fileformat.

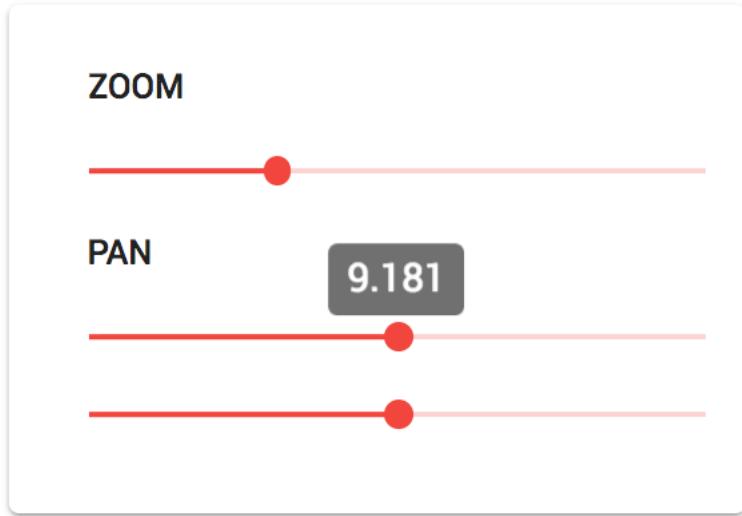


Figure 9: Zoom Component, the user is dragging the zoom slider currently, the mouse pointer is not depicted.

5.3 Base

The implementation is making a distinction between plugin and non-plugin files. The amount of non-plugin files was aimed for to be as low as possible, as they are inflexible in all output configurations and will have slightly different behaviour while developing in regard to reloads. The files not contained in plugins are the following:

- *AppState.tsx*, the mobx-state-tree root node representing the whole application state in its leafs, detailed in section 5.6.
- *BTFFFile.tsx*, containing the fileformat implementation and utility functions, described in section 5.5.
- *Hook.tsx* and *HookManager.tsx* which provide the whole dynamic interaction system between the different plugins, shown in section 5.4.
- *Loader.tsx*, *electron/index.tsx*, *renderer/index.tsx* and *web/index.tsx*, providing the loading functionality. The Electron application has two entry points, one for the main process, which is *electron/index.tsx* and one for the in-browser content, which is *renderer/index.tsx*. The in-browser one and the plain browser entry point *web/index.tsx* both call the *Loader.tsx* to initialise the state management and mount the root React component, so the user can interact finally. The Loader also handles hot-reloading, it will receive the changed source code from Webpack and update the plugins accordingly.
- *Plugin.tsx* defining the base class for a plugins, further explained in section 5.9.
- *types.d.ts* is providing the custom ambient type declarations for software dependencies, which are not providing TypeScript types on their own. In case new dependencies are added, they are likely to require and addition there.

- *Util.tsx* providing general helper functions, largely related to some math functions for texture coordinate handling.
- *loaders/glsify-loader/index.js* is the custom Webpack loader for *.glsl* files, allowing e.g. the `import zoomShader from './zoom.glsl'` statement and setting up Webpack to contain the shader source in the final bundle.
- *loaders/oxrtidatatem/OxrtiDataTextureLoader.tsx* is providing direct texture loading from in-memory BTF files, discussed in section 5.8.

5.4 Hooks

The hook system allows stable and prioritized interactions between the different plugins. All available hooks are declared inside the *Hook.tsx* file, which offers 3 different types of hooks:

```

1 // Hooks are sorted in descending priority order in their respective
2   ↳ `HookManager`
3
4 // Generic single component hook, usually used for rendering a dynamic
5   ↳ list of components
6
7 // Generic single function hook, usually used for notifications
8
9 // Generic hook config, requiring more work at the consumer side
10
11 export type HookBase = { priority?: number }

export type ComponentHook<P = PluginComponentType> = HookBase & {
  component: P }

export type FunctionHook<P = (...args: any[]) => any> = HookBase & { func:
  P }

export type ConfigHook<P = any> = HookBase & P

```

These types are used to first declare single hook types (which will be discussed within the plugins consuming them) and then construct the whole hook configuration tree for all plugins:

```

1 type HookTypes = {
2   ActionBar?: ConfigHook<ActionBar>,
3   AfterPluginLoads?: FunctionHook,
4   AppView?: ComponentHook,
5   ...
6 }

```

A plugin then can link itself into these hooks with its *hooks* method, for example:

```

1 get hooks () {
2   return {
3     // register things for the ViewerSide hook / add components to the
4       ↳ side bar

```

```

4   ViewerSide: {
5     // a plugin can register itself multiple times with different names
6     // and configurations
7     Metadata: {
8       component: BTFFMetadataConciseDisplay,
9       // hooks will be sorted internally in priority order, highest
10      // first
11      priority: -110,
12    },
13    Open: {
14      component: Upload,
15      priority: 100,
16    },
17  }

```

The state manager (section 5.6) collects all hooks and merges them into the respective HookManagers, which are then used to iterate/map over these hooks with 4 convenience functions `forEach`, `asyncForEach`, `forEachReverse`, `map` and `pick`.

5.5 BTF File

The full standalone BTF file format specification can be found inside Appendix A. The implementation in `BTFFFile.tsx` is an in-memory implementation of that file with following interface, it is mainly a ‘dumb’ data container.

```

1  export default class BTFFFile {
2    /** running id numbers to allow easy cache busts */
3    id: number;
4    /** JSON object of the included oxrti state */
5    oxrtiState: object;
6    /** default data representation */
7    data: Data;
8    /** reference to annotation layers */
9    layers: AnnotationLayer[];
10   /** user visible name */
11   name: string;
12   /** manifest can come from an unpacked zip, usually typed as any */
13   constructor(manifest?: BTFFFile);
14   /** canonical zip name for name and id */
15   zipName(): string;
16   /** return true if no data is contained/is dummy object */
17   isDefault(): boolean;
18   /** export the JSON data of the manifest.json file */
19   generateManifest(): string;
20   /** export user visible shortened metadata */
21   conciseManifest(): string;
22   /**

```

```

23      * Generate a unique tex container which the gl-react loader will
24      ← cache
25      * @param channel reference to the named channel
26      * @param coefficent reference to the named child coefficent of
27      ← channel
28      */
29      texForRender(channel: string, coefficent: string): TexForRender;
30      /**
31      * Generate a tex configuration for a layer
32      * @param id of the layer, must be found in this.layers
33      */
34      annotationTexForRender(id: string): TexForRender;
35      /** aspect ratio of the contained data */
36      aspectRatio(): number;
37      /** package the current data into a zip blob */
38      generateZip(): Promise<Blob>;
39  }
40  /** unpackage a zip blob into a BTFFfile */
41  export declare function fromZip(zipData: Blob | ArrayBuffer):
42      ← Promise<BTFFfile>;

```

Notable is the PreDownload hook which is called before the generateZip function is called, to allow the ImpExp plugin and the Paint plugin to fill the BTF with their respective data. This hook is called by the AppState though, as the BTFFfile implementation is fully standalone (apart from the jszip dependency), in case other software wishes to re-use the implementation.

5.6 State Management

The *AppState.tsx* file is the root of the applications state tree, but itself only contains limited data:

```

1 // types refering to MST types
2 const AppStateData = types.model({
3     // keep references to loaded plugins
4     plugins: types.late(() => types.optional(types.map(Plugin), {})),
5     // have a HookManager for each HookName
6     hooks: types.optional(types.map(HookManager), {}),
7     // currently opened BTF file, name is the key for the BTFCache
8     currentFile: '',
9 })

```

Its main function is to load all plugins and then let them share data between each other via the hook system. The set of plugins to be loaded is defined inside *oxrti.plugins.json*, the plugins defined there will be loaded in order. The state tree after initial plugin load is shown in Figure 11. As only `@actions` can modify the state, it is easy to follow the changing app state, as shown in Figure 10.

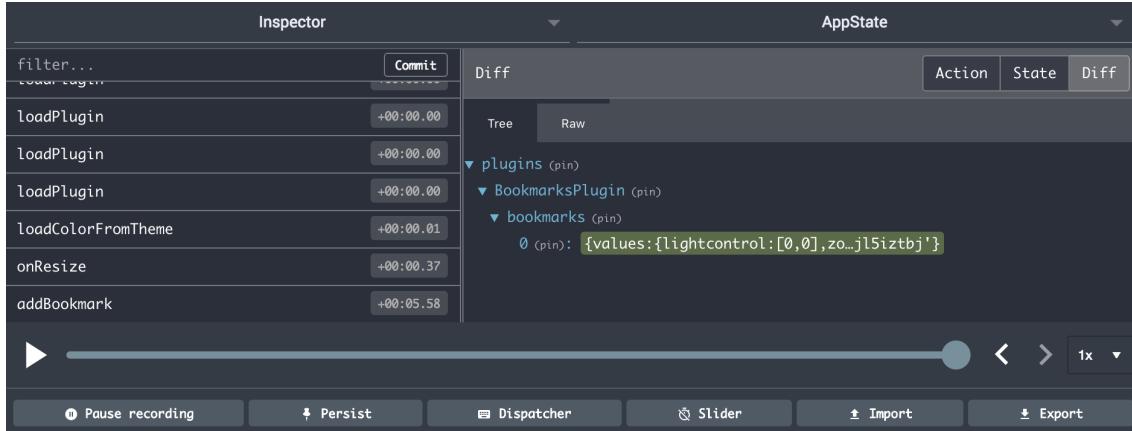


Figure 10: Run MobX at application startup and after one click on ‘Add Bookmark’. The differential of that action is on the right.

5.7 Renderer Stack

Even though the renderer stack is fully contained in the plugins, its principles span most of them, thus an overview is in order. All WebGL rendering is done through the gl-react library. If a plugin wants to register a node inside the stack, it is using the hook system, e.g. the Rotation plugin:

```

1 // register two nodes inside the rendering stack
2 // higher priority will be run first
3 ViewerRender: {
4     // first center the underlying texture
5     Centerer: {
6         component: CentererUI,
7         inversePoint: this.undoCurrentCenterer,
8         forwardPoint: this.doCurrentCenterer,
9         priority: 11,
10    },
11    // then rotate it
12    Rotation: {
13        component: RotationUI,
14        inversePoint: this.undoCurrentRotation,
15        forwardPoint: this.doCurrentRotation,
16        priority: 10,
17    },
18 },

```

The registered nodes are components again, which will link them into the automated mobx reactions, for example the Centerer node:

```

1 export const CentererUI = Component(function CentererNode (props) {
2     // dynamic sizes depending on the loaded btf
3     // if the btf changes, the uniforms will be updated automatically
4     let [width, height] = this.centererSizes
5     let maxDims = this.maxDims

```

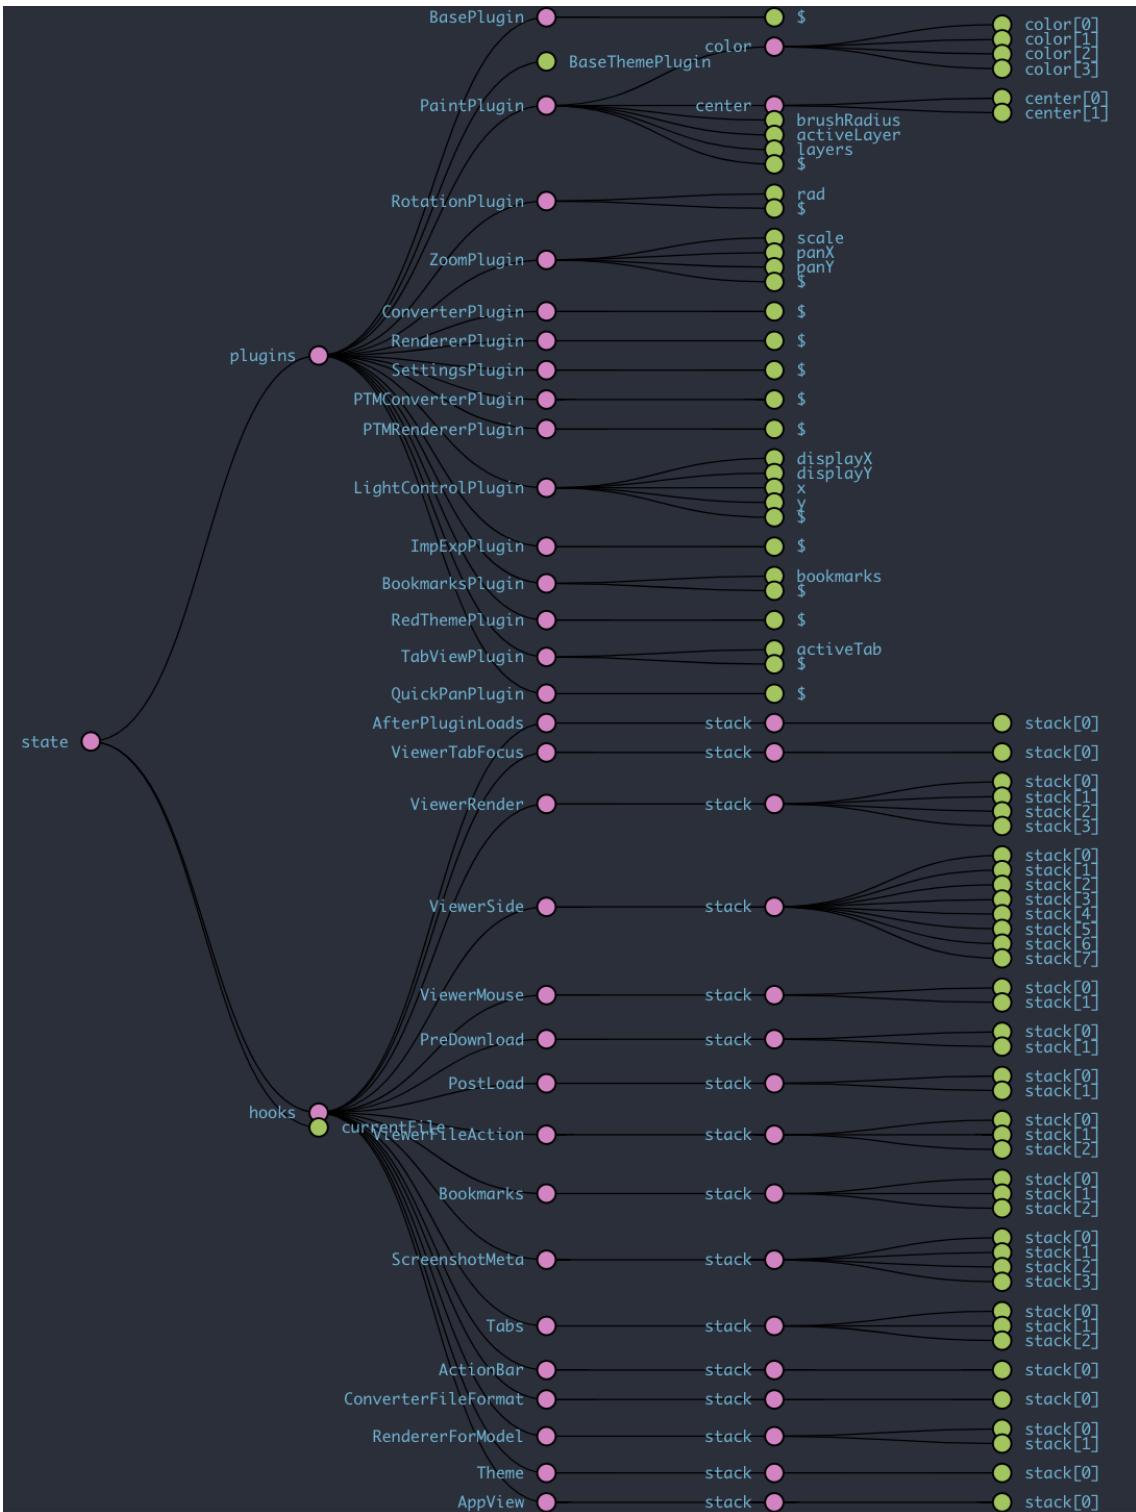


Figure 11: State tree after initial application load. All plugins are initialized and their hooks are registered. Visualized with redux dev tools.[37]

```

6   // create one gl-react Node
7   return <Node
8     width={maxDims}
9     height={maxDims}
10    shader={{{
11      // shader from import centerShader from './centerer.gsl'
12      frag: centerShader,
13    }}}
14    // uniforms will be automatically type-converted to the appropriate
15    // → WebGL types
16    uniforms={{{
17      // referring to rendering output one step before
18      children: props.children,
19      inputHeight: height,
20      inputWidth: width,
21      maxDim: maxDims,
22    }} />
23  })

```

For the currently delivered default configuration the following nodes are registered:

1. PaintNode
2. CentererNode
3. RotationNode
4. ZoomNode

The Renderer plugin will start by picking the proper base rendering node depending on the channel format of the currently loaded BTF file (btf inside the code):

```

1  props.appState.hookForEach('RendererForModel', (Hook) => {
2    if (Hook.channelModel === btf.data.channelModel) {
3      current = <Hook.node
4        key={btf.id}
5        lightPos={lightControl.lightPos}
6        renderingMode={this.renderMode}
7      />
8    }
9  })

```

This initial node will then be wrapped by the registered nodes:

```

1  props.appState.hookForEach('ViewerRender', (Hook) => {
2    current = <Hook.component
3      key={btf.id}
4      >{current}</Hook.component>
5  })

```

5.8 Texture Loader

The bridge between the loaded BTF files and the WebGL contexts needs to be closed with custom code, as gl-react was not supporting the load of PNG textures from memory. An abridged version of the code follows, as it is a good example of the Promise pattern used in some following parts.

```
1  loadTexture(config: TexForRender) {
2      // keep track of the amount of currently loading textures
3      AppState.textureIsLoading()
4      // shortcut for the WebGL reference
5      let gl = this.gl
6      // access the raw data buffer from the texture configuration
7      let data = config.data
8      // create a Promise, that basically is chaining callbacks together and
9      // then asynchronously executing each step
10     let promise =
11         // first create an ImageBitmap object, we need to flipY as the
12         // textures
13         are orientated naturally inside the BTF file but WebGL is
14         // expecting them bottom row first
15         createImageBitmap(data, { imageOrientation: 'flipY' })
16         // the catch step is called if the previous part failed
17         .catch((reason) => {
18             // some browsers do not like loading a lot of big textures in
19             // parallel and will garbage collect them in between and thus
20             // fail, as a fallback the limiter function is used to limit
21             // concurrency of that part
22             return limiter(() => createImageBitmap(data))
23         })
24         // the then part is called when the previous step succeeded
25         .then(img => {
26             // create and bind a new WebGL texture
27             let texture = gl.createTexture()
28             gl.bindTexture(gl.TEXTURE_2D, texture)
29             let type: number
30             // map the type from the BTF file to a WebGL texture type
31             switch (config.format) {
32                 case 'PNG8':
33                     type = gl.LUMINANCE
34                     break
35                     // ...
36                 case 'PNG32':
37                     type = gl.RGBA
38                     break
39             }
40             // load the imageBitmap into the texture
41             gl.texImage2D(gl.TEXTURE_2D, 0, type, type, gl.UNSIGNED_BYTE,
42             img)
```

```

36         gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_MIN_FILTER,
37             ↪ gl.NEAREST)
38         gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_MAG_FILTER,
39             ↪ gl.NEAREST)
40         // finished and return
41         AppState.textureLoaded()
42         return { texture, width: img.width, height: img.height }
43     })
44     .catch((reason) => {
45         // a null texture will be empty
46         alert('Texture failed to load' + reason)
47         AppState.textureLoaded()
48         return { texture: null, width: config.width, height:
49             ↪ config.height }
50     })
51     // the calling code will have its own catch/then logic
52     return promise
53 }

```

5.9 Plugins

All plugins extend the relatively limited abstract plugin class in *Plugin.tsx* and have no preserved state:

```

1  /** General Plugin controller, will be loaded into the MobX state tree */
2  export declare class PluginController extends PluginShim {
3      /** referential access to app state, will be set by the plugin loader
4          ↪ */
5      AppState: IAppState;
6      /** called when the plugin is initially loaded from file */
7      load(AppState: IAppState): void;
8      /** all hooks the plugin is using */
9      readonly hooks: HookConfig;
10     /** get a single typed hook */
11     hook<P extends HookName>(name: P, instance: string): HookType<P>;
12     /** called before the plugin will be deleted from the state tree,
13         ↪ ususally used for volatile state fixes, e.g. paint layers */
14     hotUnload(): void;
15     /** called after the plugin was restored in the state tree */
16     hotReload(): void;
17     /** convenience function to inverse a rendering point from surface
18         ↪ coordinates into texture coordinates */
19     inversePoint(point: Point): Point;
20     /** some components need references to their actual DOM nodes, these
21         ↪ are stored outside the plugins scope to allow hot-reloads */
22     handleRef(id: string): (ref: any) => void;
23     /** return a stored ref */
24     ref(id: string): any;

```

```
21 }
```

The most important function is the `PluginCreator`, which merges the mobx-state-tree model with the classy-mst controller code and provides a wrapper function to create components bound to the containing plugin:

```
1 /**
2  * Create Subplugins
3  * @param Code is the controller
4  * @param Data is the model
5  * @param name must be the same as the folder and filename
6 */
7 function PluginCreator<S extends ModelProperties, T, U> (Code: new () =>
8   U, Data: IModelType<S, T>, name: string) {
9   // create the resulting plugin class
10  let SubPlugin = mst(Code, Data, name)
11  // higher-order-component
12  // inner is basically (props, classes?) => ReactElement
13  // inner this will be bound to the SubPlugin instance
14  type innerType<P, C extends string> = (this: typeof SubPlugin.Type,
15    props: ComponentProps & { children?: ReactNode } & P, classes?:
16    ClassNameMap<C>) => ReactElement<any>
17  // P are they freely definable properties of the embedded react
18  // component
19  // C are the inferred class keys for styling, usually no need to
20  // manually pass them
21  function SubComponent<P = {}, C extends string = ''> (inner:
22    innerType<P, C>, styles?: StyleRulesCallback<C>):
23    PluginComponentType<P> {
24      // wrapper function to extract the corresponding plugin from props
25      // into plugin argument typedly
26      let innerMost = function (props: any) {
27        let plugin = (props.appState.plugins.get(name)) as typeof
28          SubPlugin.Type
29        // actual rendering function
30        // allow this so all code inside a plugin can just refer to
31        // this
32        let innerProps = [props]
33        // append styles
34        if (styles)
35          innerProps.push(props.classes)
36        // call the embedded component
37        return inner.apply(plugin, innerProps)
38      };
39      // set a nice name for the MobX/redux dev tools
40      (innerMost as any).displayName = inner.name
41      // use MobX higher order functions to link into the state tree
42      let func: any = inject('appState')(observer(innerMost))
43      // wrap with material-ui styles if provided
44      if (styles)
```

```

35         func = withStyles(styles)(func);
36         // also name the wrapped function for dev tools
37         (func as PluginComponentType<P>).displayName =
38             → `PluginComponent(${inner.name})`
39         return func
40     }
41     // allow easier renaming in the calling module
42     return { Plugin: SubPlugin, Component: SubComponent }
43 }

```

A minimal example would be:

```

1  const BasePluginModel = Plugin.props({
2     greeting: 'In the beginning was the deed!',
3 })
4
5  class BasePluginController extends shim(BasePluginModel, Plugin) {
6     @action
7     onGreeting (event: any) {
8         this.greeting += '!'
9     }
10 }
11
12 // general plugin template code
13 const { Plugin: BasePlugin, Component } =
14     → PluginCreator(BasePluginController, BasePluginModel, 'BasePlugin')
15 export default BasePlugin
16 // export the type to allow other plugins to retrieve this plugin
17 export type IBasePlugin = typeof BasePlugin.Type
18
19 // CSS styles, classnames will be mangled, so styles is passed to the
20 // component
21 const styles = (theme: Theme) => createStyles({
22     hello: {
23         color: 'red',
24     },
25 })
26
27 // props are standard react props, classes contains the mangled names
28 const HelloWorld = Component(function HelloWorld (props, classes) {
29     return <p className={classes.hello}
30         → onClick={this.onGreeting}>{this.greeting}</p>
31 }, styles)

```

The rest of the plugins are presented on a higher conceptual level, as their internal APIs are most times used only by themselves. Their hooks will be discussed though.

5.9.1 Base Plugin

The base plugin is containing no further internal logic, but only is providing some shared display components. These could theoretically have been implemented outside of any plugin, but putting them into the Base plugin simplifies the distinction between components bound to plugins and unbound plugins, by having no unbound plugins at all. It also simplifies the code loading, as just the Base plugin can be (re-)loaded like all other plugins. The provided components are:

- JSONDisplay, showing a JSON object in a prettified form, used for displaying different metadata objects
- BTFFMetadataDisplay, showing the currently loaded BTF's data
- RenderHooks, to render all components attached to a hook
- SafeGLIInspector, wrapping the secondary WebGL surfaces and disabling theme, if WebGL debug tools are used, which can handle only one surface.
- Tooltip, wrapping the material-ui tooltip component to fix styling errors

5.9.2 BaseTheme, RedTheme and BlueTheme Plugin

The BaseTheme plugin is having two essential properties: `themeBase` and `themeExtension`. `themeBase` contains the app wide configuration for the MaterialUI library. Concrete themes (at the moment RedTheme and BlueTheme plugins) extend this BaseTheme and then set their `themeExtension` according to their modifications. To pick a theme to be used the ThemeConfig hook exists:

```
1 type ThemeConfig = {
2     controller: { theme: Theme },
3 }
```

Concrete themes register with:

```
1 get hooks (): HookConfig {
2     return {
3         Theme: {
4             Red: {
5                 priority: 100,
6                 controller: this,
7             },
8         },
9     }
10 }
```

If multiple plugins are registering themes, the appState will pick the theme with the highest priority to apply to the app.

5.9.3 TabView Plugin

The TabView plugin is providing the full app experience and is targeted at the Electron output and the online hosted version. As a tabbed container will occasionally delete the content of non-active tabs, some hooks are needed to ensure graceful behaviour of all displayed tabs:

```
1 // register a new tab
2 type Tab = {
3     // component to be the base of the tab
4     content: PluginComponentType
5     tab: TabProps,
6     padding?: number,
7     // async functions to allow customisation before/after tabs change
8     beforeFocusGain?: () => Promise<void>,
9     afterFocusGain?: () => Promise<void>,
10    beforeFocusLose?: () => Promise<void>,
11    afterFocusLose?: () => Promise<void>,
12 }
13
14 // action buttons on the top rights
15 type ActionBar = {
16     onClick: () => void,
17     title: string,
18     enabled: () => boolean,
19     tooltip?: string,
20 }
21
22 // notifications if the tab changes for sub-components which are not being
23 // → a tab themselves
23 type ViewerTabFocus = {
24     beforeGain?: () => void,
25     beforeLose?: () => void,
26 }
```

This view is depicted in Figure 12.

5.9.4 SingleView Plugin

The SingleView plugin is aimed to provide a contained viewer experience, e.g. on a museum's website, it is just displaying the tab with the highest priority, see Figure 12.

5.9.5 Converter Plugin

The Converter plugin consists of multiple parts:

- The converter user interface, as shown in Figure 13.

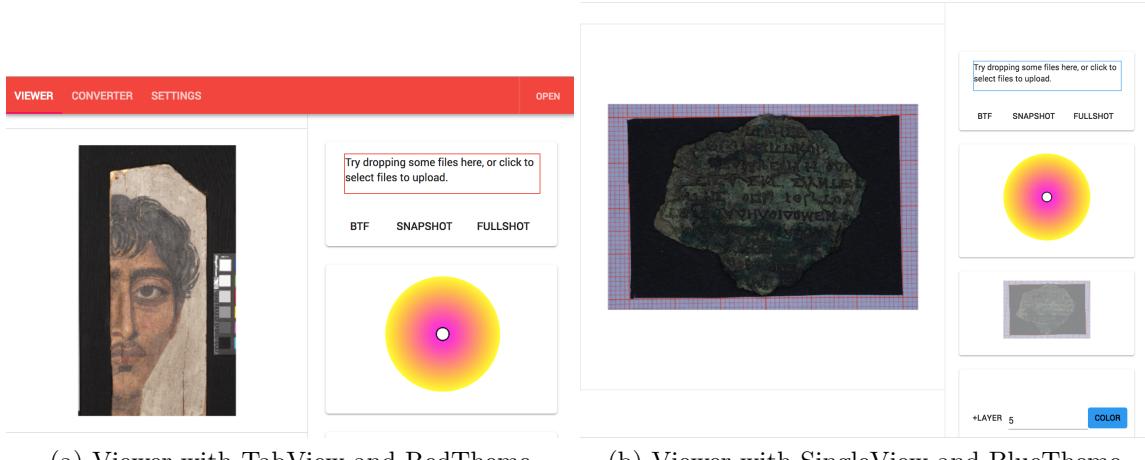


Figure 12: Only changes are two lines in *oxrти.plugins.json*.

- BMPWriter and PNGWriter, both extending Writer to write the converted textures. The BMPWriter is customised as no current package is offering the required functionality. The PNGWriter is wrapping pngjs.
- An abstract base converter strategy, to be extended by plugins providing a concrete converter with the following interface:

```

1  export default abstract class ConverterStrategy {
2    /** raw file buffer */
3    fileBuffer: ArrayBuffer;
4    /** wrapped file buffer */
5    inputBuffer: Buffer;
6    /** UI delegate for status updates */
7    ui: IConverterUI;
8    /** pixelData buffer, pointing into fileBuffer+metadata length */
9    pixelData: Buffer;
10   /** extracted width and height */
11   width: number;
12   height: number;
13   /** concrete BTF output */
14   output: BTFFile;
15   /** freeform format depending JSON object, e.g. biases for PTMs */
16   formatMetadata: object;
17   /** current channelModel as defined in the BTF specification */
18   channelModel: ChannelModel;
19   /** total pixel count */
20   readonly pixels: number;
21   /** started from the converter ui */
22   constructor(content: ArrayBuffer, ui: IConverterUI);
23   /** pointer into the raw file buffer */
24   currentIndex: number;
25   /** read metadata till newline */
26   readTillNewLine(): string;

```

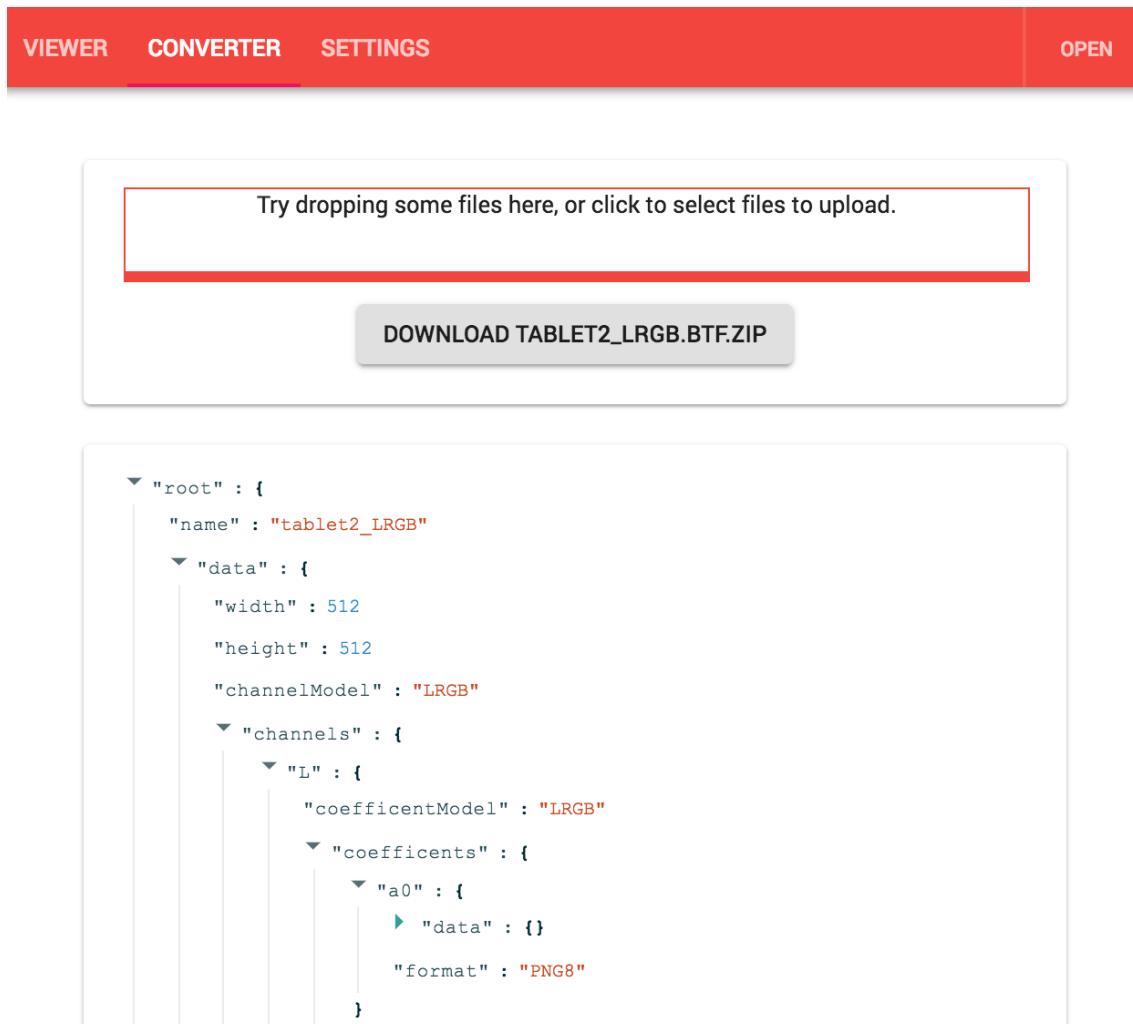


Figure 13: Converter user interface, with display of the extract manifest.

```

27     /** read one item, usually byte */
28     readOne(): number;
29     /** prepare pixeldata buffer */
30     preparePixelData(): Promise<void>;
31     /** run the actual conversion */
32     process(): Promise<BTFFFile>;
33     /** only these need to implemented by concrete converters */
34     /** read the metadata block */
35     abstract parseMetadata(): Promise<void>;
36     /** read the pixel block */
37     abstract readPixels(): Promise<void>;
38     /** read potential suffix after pixel block */
39     abstract readSuffix(): Promise<void>;
40     /** bundle the read pixels into channels according to the channelModel
41      ↵ */
41     abstract bundleChannels(): Promise<Channels>;
42 }

```

If a plugin wants to register a concrete converter it would use following hook:

```

1 ConverterFileFormat: {
2   PTM: {
3     fileEndings: ['.ptm'],
4     // referring to the class extending the base strategy
5     strategy: PTMConverterStrategy,
6   },
7 }

```

5.9.6 PTMConverter Plugin

The PTM Converter plugin is converting *.ptm* files, as they are described in section 2.2. Currently the RGB and LRGB light models are supported. Their main interesting part is the pixel data reader:

```

1 async readPixelsRGB () {
2   // allocate output buffers for ['R0R1R2', 'R3R4R5', 'G0G1G2',
3   //   ↵ 'G3G4G5', 'B0B1B2', 'B3B4B5']
4   this.coeffData = this.coeffNames.map(e => Buffer.alloc(this.pixels *
5   //   ↵ 3))
6   // this.pixelData contains a block of pixels * [a_0, a_1, a_2, a_3,
7   //   ↵ a_4, a_5] for each color
8   for (let y = 0; y < this.height; ++y) {
9     for (let x = 0; x < this.width; ++x) {
10       for (let color = 0; color <= 2; color++) {
11         // iterate over pixels and colors
12         let inputIndex = (((y * this.width) + x) + this.pixels *
13         //   ↵ color) * 6
14         // flip the Y index, as PTM is bottom-row first, whereas
15         //   ↵ BTF is top-row first

```

```

11     let targetIndex = (((this.height - 1 - y) * this.width) +
12         ↪ x) * 3
13     // iterate over the coefficents for the given
14     ↪ pixel/inputIndex */
15     for (let i = 0; i <= 5; i++) {
16         let bucket = color * 2 + Math.floor(i / 3)
17         this.coeffData[bucket][targetIndex + (i % 3)] =
18             ↪ this.pixelData[inputIndex + i]
19     }
20     // update the progress bar
21     if (y % (this.height / 100) === 0) {
22         await this.ui.setProgress(y / this.height * 100 + 1)
23     }
24 }

```

5.9.7 Renderer Plugin

The Renderer plugin is providing the main user interface, which is split into two parts, on the left the rendered object and on the right further controls. As such it is using a comprehensive set of hooks:

```

1 // specific channelModel renderers can register their base node
2 type BaseNodeConfig = {
3     channelModel: ChannelModel,
4     // a basenode can declare that it supports multiple rendering modes
4     ↪ e.g. default, surface normals, specular enhancement, etc.
5     renderingModes: string[],
6     node: PluginComponentType<BaseNodeProps>,
7 }
8
9 // hook for a node in the renderer stack
10 type RendererNode = {
11     component: PluginComponentType,
12     // if the node is transforming the texture coordinates, and inverse
12     ↪ method needs to be provided
13     inversePoint?: (point: Point) => Point,
14     forwardPoint?: (point: Point) => Point,
15 }
16
17 // hook for listening to mouse event inside the main renderer
18 type MouseConfig = {
19     dragger: (oldTex: Point, nextTex: Point, oldScreen: Point, nextScreen:
19     ↪ Point, dragging: boolean) => boolean,
20     mouseUp?: (nextScreen: Point, nextTex: Point) => boolean,
21     mouseLeft?: () => void,
22 }

```

```

23
24 // hook for adding file actions/buttons below the upload field
25 type ViewerFileAction = {
26     tooltip: string,
27     text: string,
28     action: () => Promise<void>,
29 }
30
31 // hook for adding information to the metadata file when shots are
32 // exported
32 type ScreenshotMeta = {
33     key: string,
34     fullshot?: () => (string | number)[] | string | number,
35     snapshot?: () => (string | number)[] | string | number,
36 }
37
38 // components to be rendered inside the drawer
39 type ViewerSide = ComponentHook
40
41 // notifications, that a btf file will be exported and plugins should
42 // update their respective data inside the current in-memory version
42 type PreDownload = FunctionHook
43
44 // notification that a btf file was loaded, plugins can import extra data
45 type PostLoad = FunctionHook
46
47 // components to be rendered around the surface
48 type ViewerSurfaceAttachment = ComponentHook

```

The actual object rendering is done by the stack as described in section 5.7, the Renderer plugin is only providing a dynamically resized and centered surface for the stack to be drawn in. The surface is always kept square, even if the loaded BTF is not, to streamline and simplify texture coordinate handling.

5.9.8 PTMRenderer Plugin

The PTMRenderer Plugin is rendering the RGB and LRGB channel models. Here only the RGB is covered, as the principles for LRGB are the same. Most channel renderers will be split in two parts, one node for the rendering stack:

```

1 const coeffs = ['a0a1a2', 'a3a4a5']
2 // return a texture configuration array for the given coefficient
3 function mapper (btf: BTFFile, name: string) {
4     return coeffs.map(c => {
5         return btf.texForRender(name, c)
6     })
7 }
8
9 // render a RGB object

```

```

10 const PTMRGB = Component<BaseNodeProps>(function PTMRGB (props) {
11     let btf = props.appState.btf()
12     return <Node
13         // from ./ptmrgb.gsls
14         shader={shaders.ptmrgb}
15         // adaptive sizing if wanted
16         width={props.width || btf.data.width}
17         height={props.height || btf.data.height}
18         uniforms={{
19             // usually coming from the lightcontrol plugin, is [x:number,
20             // → y:number, z:number]
21             lightPosition: props.lightPos,
22             // texture arrays
23             texR: mapper(btf, 'R'),
24             texG: mapper(btf, 'G'),
25             texB: mapper(btf, 'B'),
26             // retrieve the untyped formatExtra
27             biases: (btf.data.formatExtra as PTMFormatMetadata).biases,
28             scales: (btf.data.formatExtra as PTMFormatMetadata).scales,
29         }} />
30     )

```

And one shader, implementing the light model described in section 2.2.

```

1 precision highp float;
2 varying vec2 uv;
3 // higher and lower coefficents per color
4 uniform sampler2D texR[2];
5 uniform sampler2D texG[2];
6 uniform sampler2D texB[2];
7 uniform float biases[6];
8 uniform float scales[6];
9 uniform vec3 lightPosition;
10
11 float channelLum(sampler2D[2] coeffsTexs, vec3 toLight) {
12     // would be unrolled by the GLSL compiler anyway
13     float a0 = texture2D(coeffsTexs[0], uv).x;
14     float a1 = texture2D(coeffsTexs[0], uv).y;
15     float a2 = texture2D(coeffsTexs[0], uv).z;
16     float a3 = texture2D(coeffsTexs[1], uv).x;
17     float a4 = texture2D(coeffsTexs[1], uv).y;
18     float a5 = texture2D(coeffsTexs[1], uv).z;
19
20     a0 = (a0 * 255.0 - biases[0]) * scales[0];
21     a1 = (a1 * 255.0 - biases[1]) * scales[1];
22     a2 = (a2 * 255.0 - biases[2]) * scales[2];
23     a3 = (a3 * 255.0 - biases[3]) * scales[3];
24     a4 = (a4 * 255.0 - biases[4]) * scales[4];
25     a5 = (a5 * 255.0 - biases[5]) * scales[5];
26

```

```

27     float Lu = toLight.x;
28     float Lv = toLight.y;
29
30     float lum = (
31         a0 * Lu * Lu +
32         a1 * Lv * Lv +
33         a2 * Lu * Lv +
34         a3 * Lu +
35         a4 * Lv +
36         a5
37     )/255.0;
38     return lum;
39 }
40
41 void main() {
42     // spotlight behaviour at the moment
43     vec3 pointPos = vec3(0,0,0);
44     vec3 toLight = normalize(lightPosition - pointPos);
45
46     float R = channelLum(texR, toLight);
47     float G = channelLum(texG, toLight);
48     float B = channelLum(texB, toLight);
49
50     gl_FragColor = vec4(R,G,B,1.0);
51 }
```

In addition to the default rendering mode the PTMRenderer plugin is also supporting a Normals rendering mode for LRGB ptms. See Figure 14 for a visual example. It is based on the proposal by MacDonald and Robson[27] and currently supports 4 extra rendering modes: X-Normal, Y-Normal, Z-Normal and False-Color-Normal. The first three render the respective value of the normal of the pixel as grey scale, the false color one maps X to R, Y to G and Z to B.

5.9.9 LightControl Plugin

The LightControl plugin is providing a visual way to control the position of the light source. The user interface is shown in Figure 15. The linear xy coordinates are transformed into hemispherical coordinates, which are then passed onto the current Base node from the Renderer plugin.

5.9.10 Rotation Plugin

The Rotation plugin allows the rotation of the viewed object (and other lower rendering nodes). It does so by first centering the object inside a large square node of the maximal spanning length when rotated ($width * Math.cos(Math.PI/4) + height * Math.sin(Math.PI/4)$), so no part will be cut off for the rendering layers on top, and then applying a simple rotation matrix on the texture coordinates on

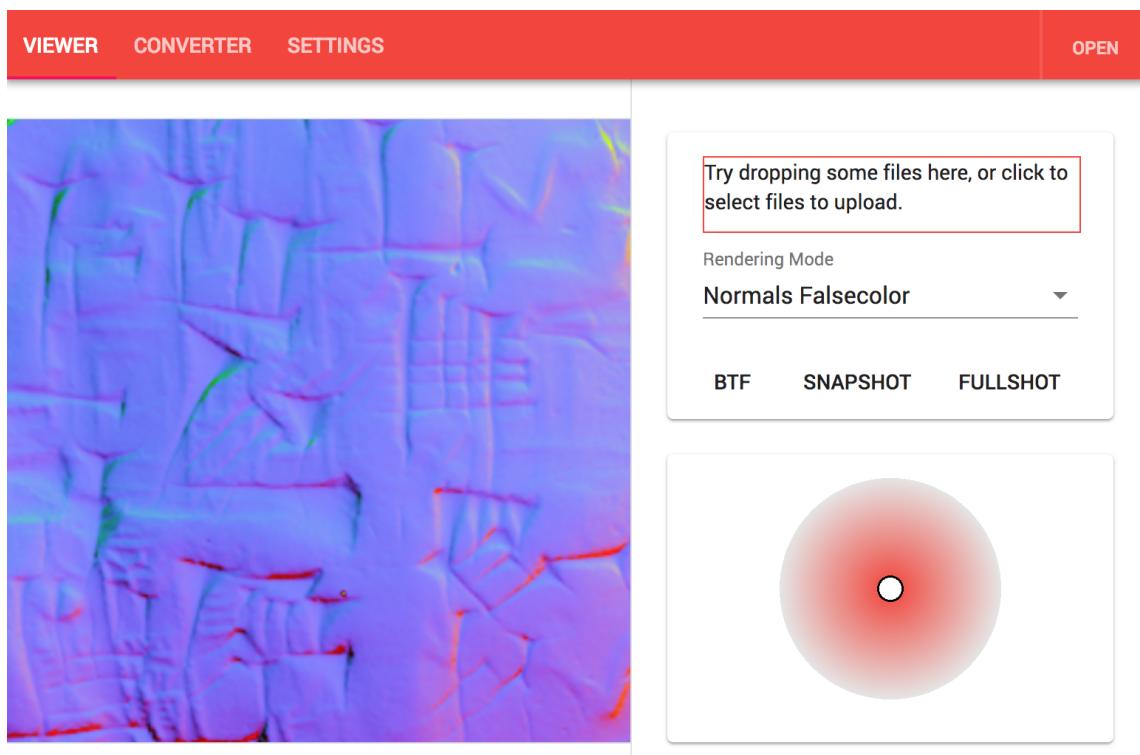


Figure 14: The Tablet ptm rendered with false color normals. The user can change the mode with the ‘Rendering Mode’ drop down on the top right.

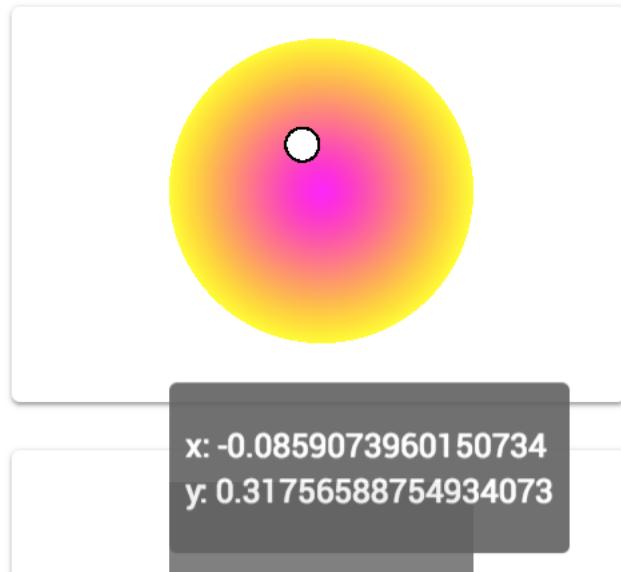


Figure 15: The user can drag around on the imaginary dome and change the light position.

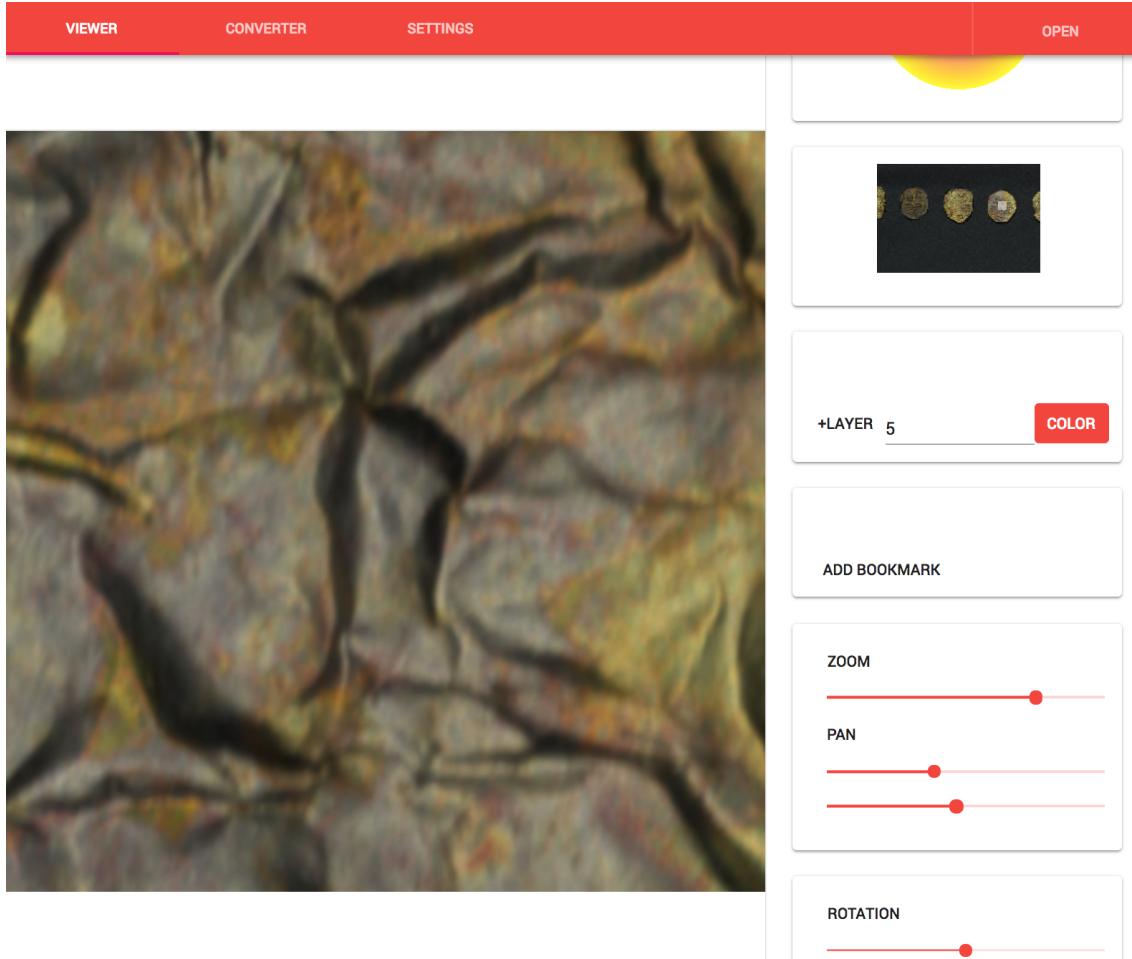


Figure 16: Applied zoom, note the small part of the actual object being shown. The shader architecture allows for seamless and performant zooming between nearest viewpoints and more removed ones.

the next node. Even though the whole renderer is rotated, the use of rotation slider will dynamically adapt the pan values to keep the previous center centered.

5.9.11 Zoom Plugin

The Zoom plugin is providing zoom and pan functionality (pan would not be necessary without zoom), an example is Figure 16. All zooming and panning is implemented as a special shader in the render stack.

5.9.12 QuickPan Plugin

The QuickPan plugin is responsible for rendering the small view box on the top right inside Figure 16. It is doing so by rendering injecting a temporary node inside the rendering stack, waiting for an object's first render and then preserving that texture in a lower resolution, by setting the `width` and `height` properties of the nodes

(compare the component in section 5.9.8). It then uses the RectRender component to render the semi-transparent rectangle of the current visible part over the captured texture. To do so, it is using the `inversePoint` function on $[0, 0]$, $[0, 1]$, $[1, 1]$, $[0, 1]$ to transform the bounds of the currently rendered (main) surface into texture coordinates of the object. With these coordinates a simple point-in-rectangle check can be done to highlight the desired area.

5.9.13 Paint Plugin

The Paint plugin adds the functionality to have overlays on the object to allow for annotations in a research context. All painting is happening as part of the renderer stack, with one node for each layer and one mixer node to combine these with the object. One interesting point is the required recompilation of the mixer shader, as the WebGL compiler is not supporting loops with unfixed amounts of maximal iterations. So for each new layer the shader source is changed. The user interface is shown in Figure 17. Up to 15 layers are supported currently, as the max texture limit is usually 16.

5.9.14 Bookmarks Plugin

The Bookmarks plugin allows bookmarks to be set for light and view configurations. Bookmarks are controlled via the following hook:

```

1 type BookmarkSaver = {
2   // key inside the bookmarks config
3   key: string,
4   // called when a new bookmark is created, any returned string or
5   // → number combination is stored
6   save: () => (string | number)[],
7   // called when a bookmark ought to be restored, the saved combination
8   // → is passed on
9   restore: (values: (string | number)[]) => void,
10 }

```

5.9.15 Notes Plugin

The Notes plugin implements notes, which are attached to a specific point of the viewed object. Each note consists of two parts: One indicator which is linked to the texture and whose position is transformed like the texture and one popup, which calculates its absolute screen coordinates to show the note's name and text. See Figure 19 for an example.

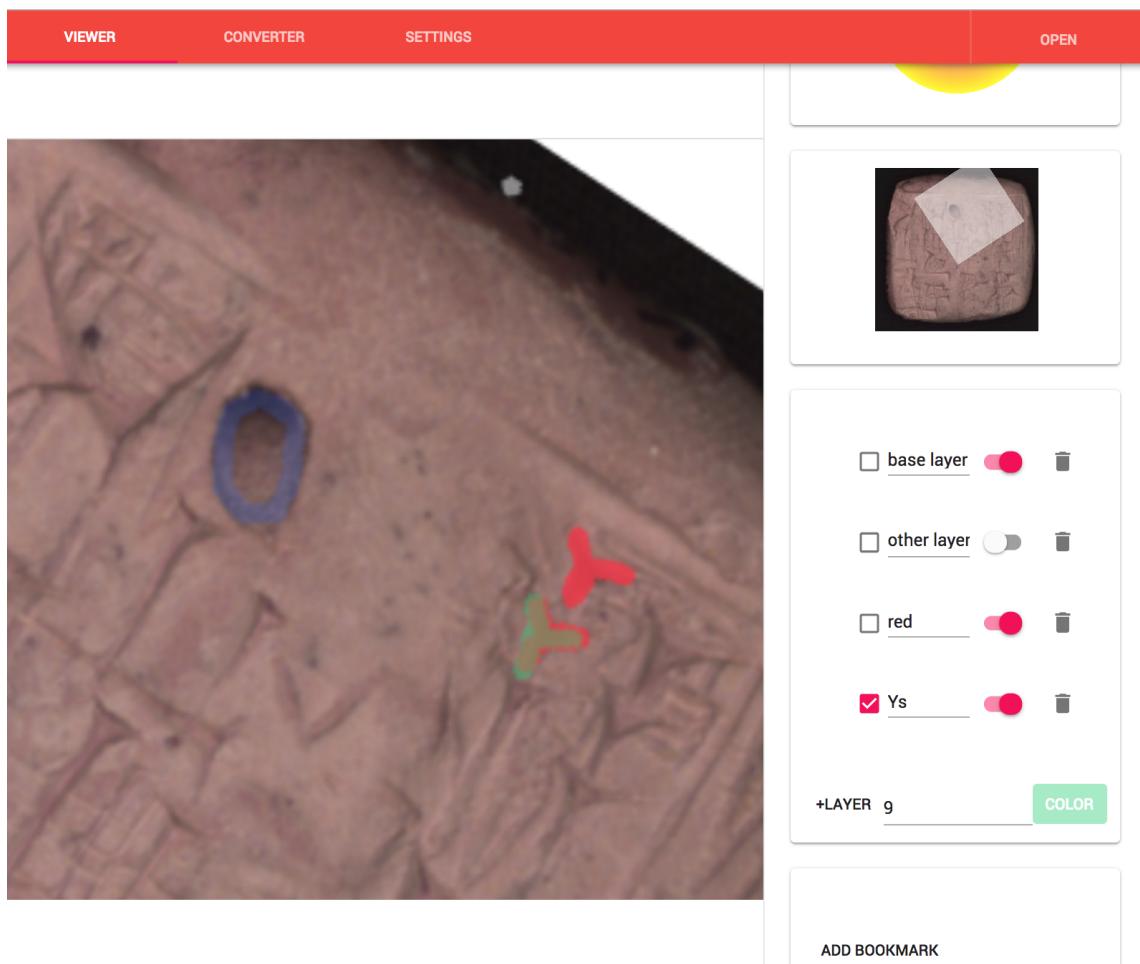


Figure 17: Paint plugin user interface. The check mark in front indicates the layer is currently drawn on. Names are freely changeable. The toggle toggles visibility of the layer. The trash icon deletes the corresponding layer. +Layer adds another layer on top. The number is the brush size. Clicking color shows a color picker.

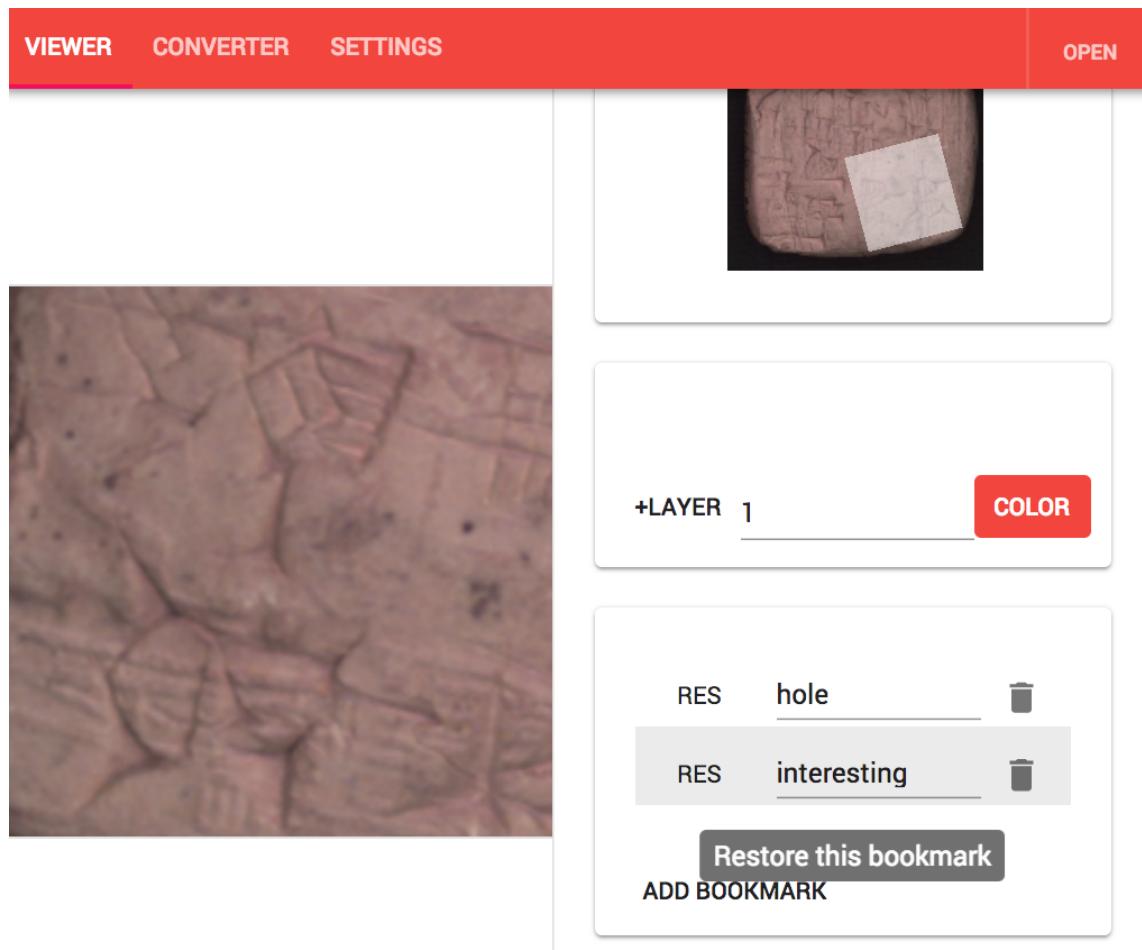


Figure 18: A click on ‘Add Bookmark’ will preserve the current light and camera positions and create an unnamed entry in the bookmark list. Clicking ‘RES’ restores this configuration. They are automatically exported when saving the current BTF.

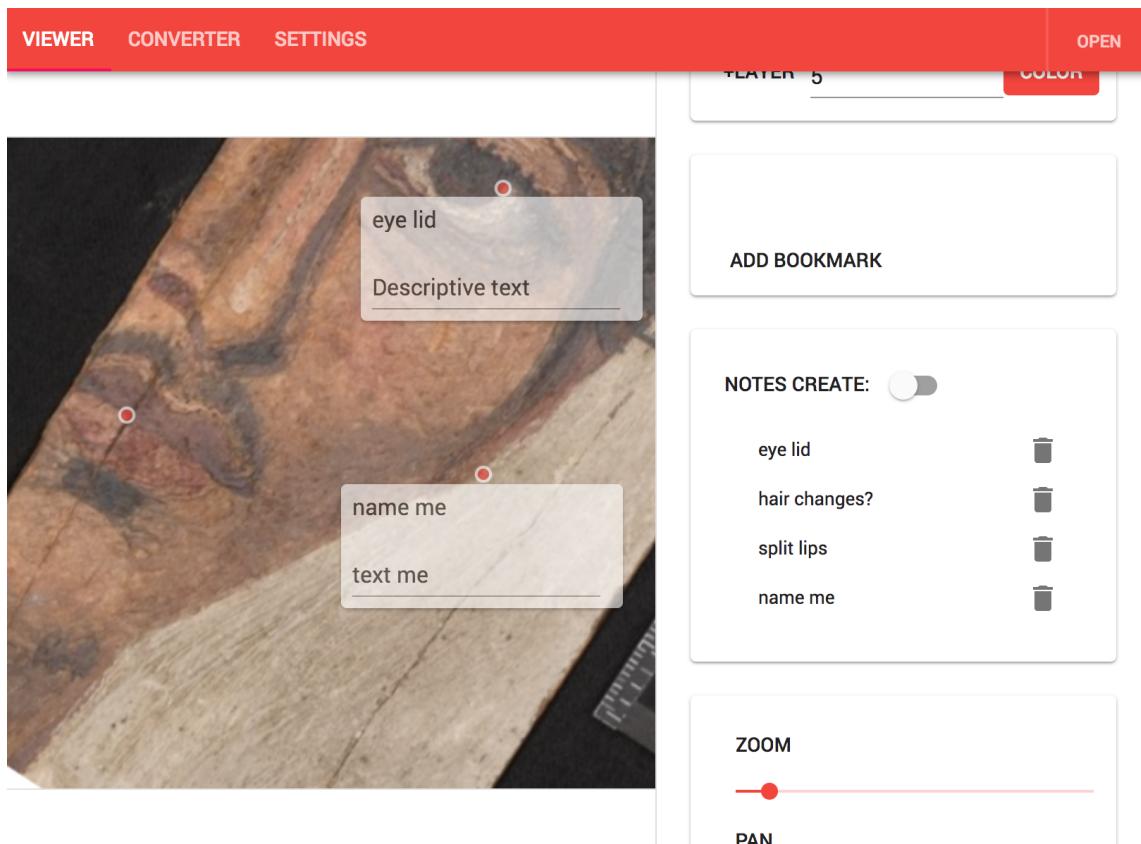


Figure 19: Notes are created by switching the ‘Notes Create:’ toggle and then clicking anywhere on the texture. The red and white dots indicate the presence of a note at that point, clicking it will toggle the display of the note. Clicking the note’s name on the right side, will center that note’s indicator.

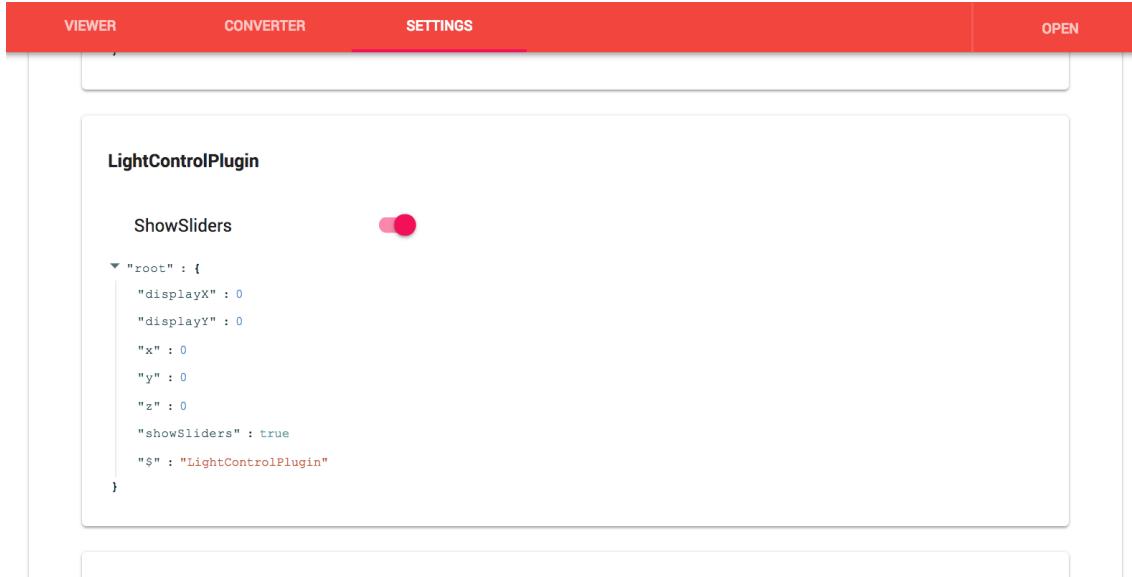


Figure 20: The LightControl plugin registers a toggle for enabling/disabling a slider based light control.

5.9.16 ImpExp Plugin

The ImpExp plugin is implementing the import and export of the current application state into/from a BTF file. For export the whole state tree is just exported as:

```
1 btf.oxrtiState = (this.appState as any).toJSON()
```

And for import the state is simply assigned as

```
1 for (let key in snapshot)
2   (this.appState as any)[key] = snapshot[key]
```

The simpleness of this plugin is a big payoff for using mobx-state-tree and classy-mst. mobx-state-tree wil take the partial snapshots and classy-mst will turn these into fresh instances of the respective plugins. All future plugins will be automatically supported, as long as they are using their models in the default way.

5.9.17 Settings Plugin

The Settings plugin provides the user with introspection into the currently loaded plugins and their configuration options. Plugins use the the SettingsConfig hook to register configuration values:

```
1 type SettingsConfig = ComponentHook & {
2   title: string,
3 }
```

The Settings plugin registers a tab and will show the configuration components, see Figure 20.

5.10 Targets

The multiple targets are configured via two mechanisms: Webpack configuration files, which are configuring the compilation and bundling of the source code and other resources. And the *oxrti.plugins.json*, which has following form:

```
1  {
2      "enabled": [
3          "BasePlugin",
4          "BaseThemePlugin",
5          // ...
6          "ZoomPlugin"
7      ],
8      "disabled": [
9          "SingleViewPlugin",
10         "UndoPlugin",
11         "TestPlugin",
12         // ...
13     ]
14 }
```

The disabled block is not used within the implementation, but it is convenient for a configurator to see which plugins could be included. The loader is then only loading the configured plugins when starting the app.

5.10.1 Electron

Most of the electron compilation is controlled by the electron-webpack package, with additions in *webpack.renderer.additions.js* and *webpack.renderer.shared.js*. The main process entry point is *src/electron/index.tsx*, which contains no further modification at this point, but could be extended in the future for more native application features. The web process' entry point is *src/renderer/index.tsx*, which is installing the electron extensions and the referring to the Loader for the rest of the application startup. Compilation is started by executing `npm run-script electronbuild` on the command line, the resulting artifacts will be in the *dist* folder, in the MacOS use case it will be *dist/oxrti-*.dmg*. The development version can be started with `npm start` on the CLI. A visual impressions is shown in Figure 21.

5.10.2 Web

The pure web target is controlled by *webpack.config.js* and *webpack.renderer.shared.js*. `npm run-script startweb` will start the local development server listening on `http://localhost:3000`. `npm run-script build` will build it statically, with *index.html* and *dist/bundle.js* referring to the latest built versions. A zip containing both files with fixed relative links is also automatically generated under *dist/oxrti.zip*. This zip can be transferred to any computer and a modern web browser should be able to run the contained implementation.

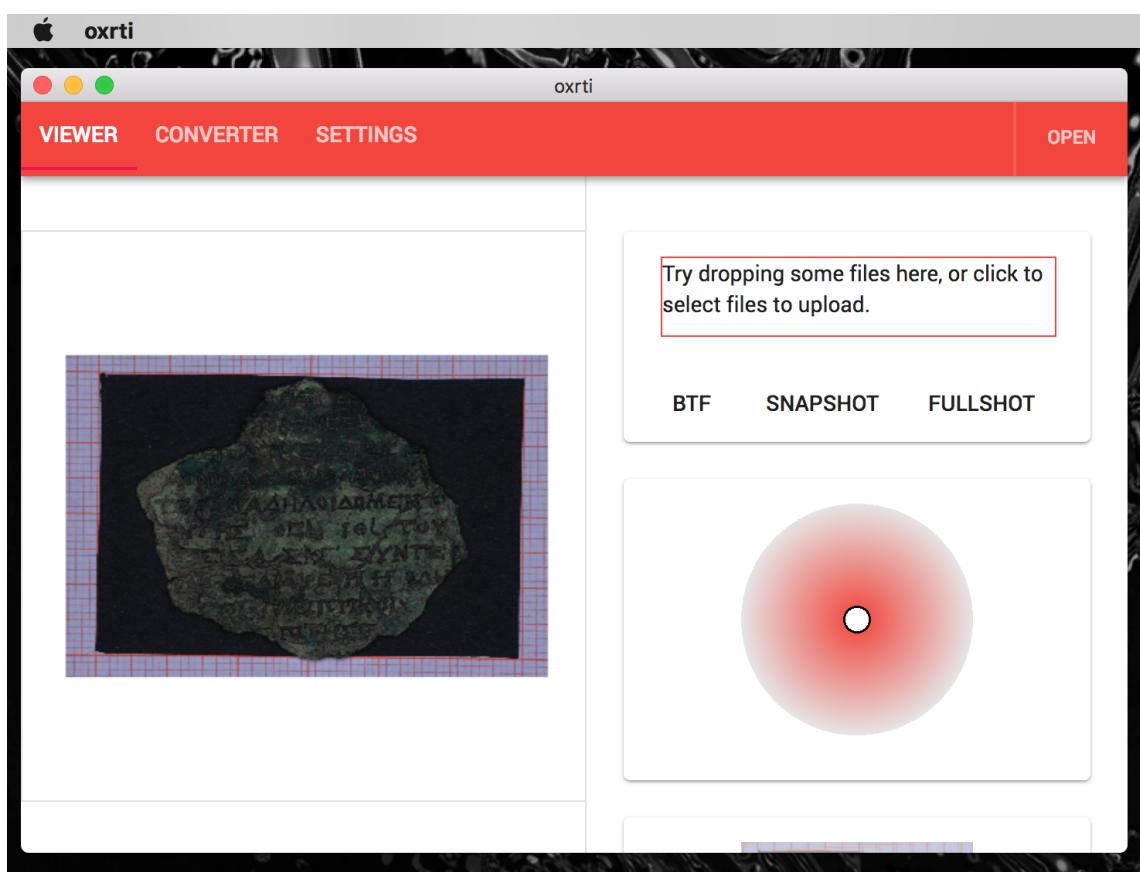


Figure 21: The bundled Electron app. Later versions will add some kind of native menu.

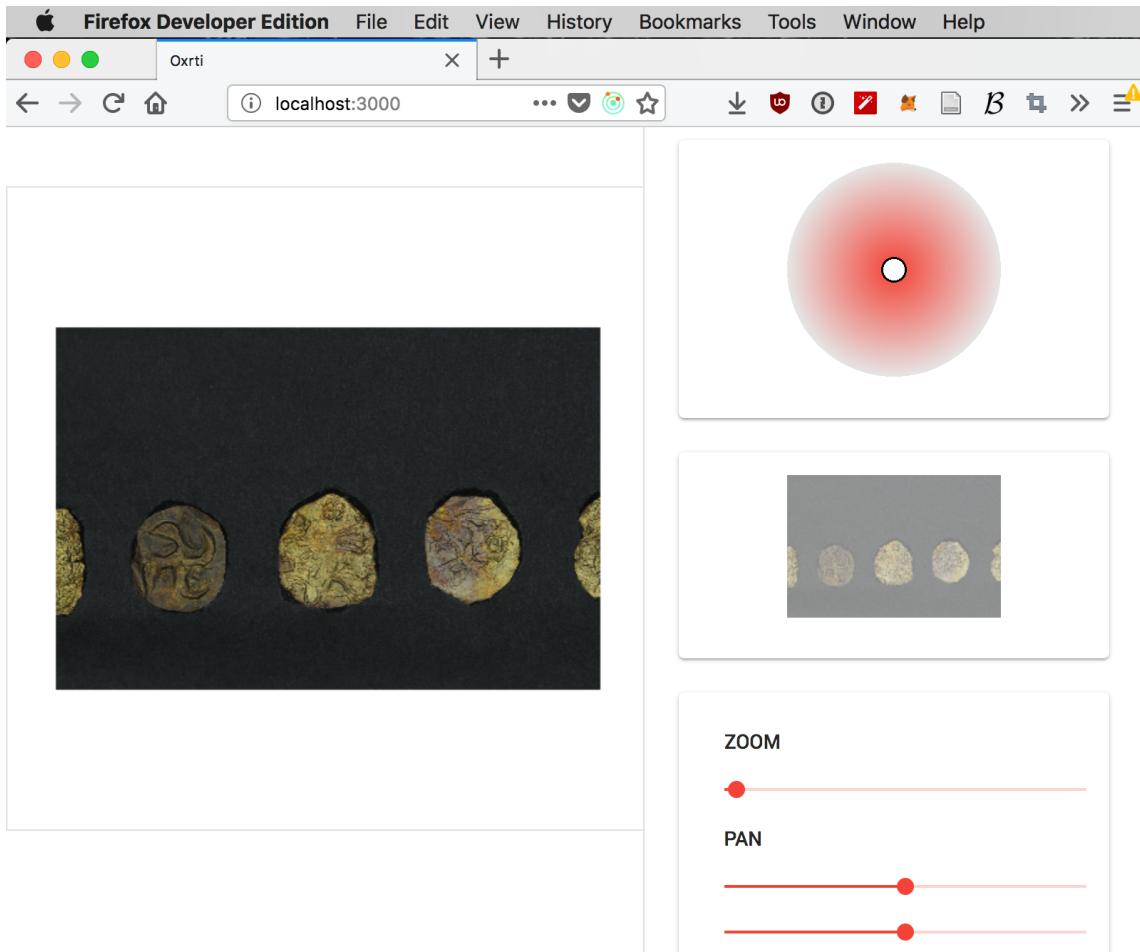


Figure 22: Embedded web version, with SingleView plugin and without Converter, Bookmarks and Paint.

5.10.3 Hosted

The build scripts are also generating a hosted version on each commit. On push the git repository[1] is automatically feeding this to an Azure Functions instance, reachable at <https://oxrtimaster.azurewebsites.net/api/azurestatic>. This process is controlled by *webpack.functions.js* and the files in *azurejs* and *azurestatic*, which are basically creating two endpoints on the azure site, each hosting the compiled *.js* file or *.html* file.

5.10.4 Embeddable

There is currently no provisioning for a more automated handling of multiple *oxrti.plugin.json* files, so it needs to be exchanged manually before building. A proposed embedded configuration for e.g. object galleries is provided by *oxrti.plugins.embedded.json*. The result is shown in Figure 22.

5.10.5 Mobile Phones

Mobile phones are supported, but currently might have problems with the distinction between *touch* and *click* events and currently have no specialized layout.

6 Results

This section summarises the results of the project and adds some quantitative analysis of the implementation. The feature set is fully reflected inside the implementation section (or the table of contents), so no additional summarization is done here. Data *file names* refer to the files distributed as [2].

6.0.6 Performance

Currently Cultural Heritage Imaging provides the most frequently used RTI viewer[13]. As a result, the performance comparisons have been made with that implementation as the reference. The data is shown in Table 1. The BTF files are about 50% the size of the ptm files, and the larger the ptm file, the better the compression benefits are. The same pattern is evident in the load time as Figure 24 shows. The bigger the files are, the more profitable the shader implementation, especially after the one-time conversion is done.

File	Pixel	Conv	Load	Load RTI	ptm size	btf size
<i>AK1A_LRGB</i>	3008x2000	4906	1713	11800	54.1	30
<i>Coin_LRGB</i>	7360x4912	24674	8224	59200	325.4	136
<i>Mummy_RGB</i>	583x1000	1769	627	5200	10.5	5.6
<i>tablet2_LRGB</i>	512x512	732	462	1600	2.4	1.2
<i>Warrior_RGB</i>	2299x3200	10590	2184	25500	132.4	50.2

Table 1: *Pixel* are the dimensions as width x height. *Conv* is the time taken by the implementation to convert the ptm file to the btf format. *Load* is the time taken by the oxrti implementation from receiving the btf file to do a complete render. *Load RTI* is the approximate time taken by the RTIViewer from opening the file to finishing the loading. Time was stopped by screenrecording, as no performance measurements are provided. The time for the processing of mipmaps and normal maps was excluded, as the oxrti implementation is not doing these steps. The time to a full first render would be even longer. *ptm size* is the file size in megabytes. *btf size* is the file size of a plain BTF file (no extra oxrti state and/or layers). Test runs done on a “MacBook Pro (Retina, 13-inch, Early 2015); 2.9 GHz Intel Core i5; 16 GB 1867 MHz DDR3; Intel Iris Graphics 6100 1536 MB”.

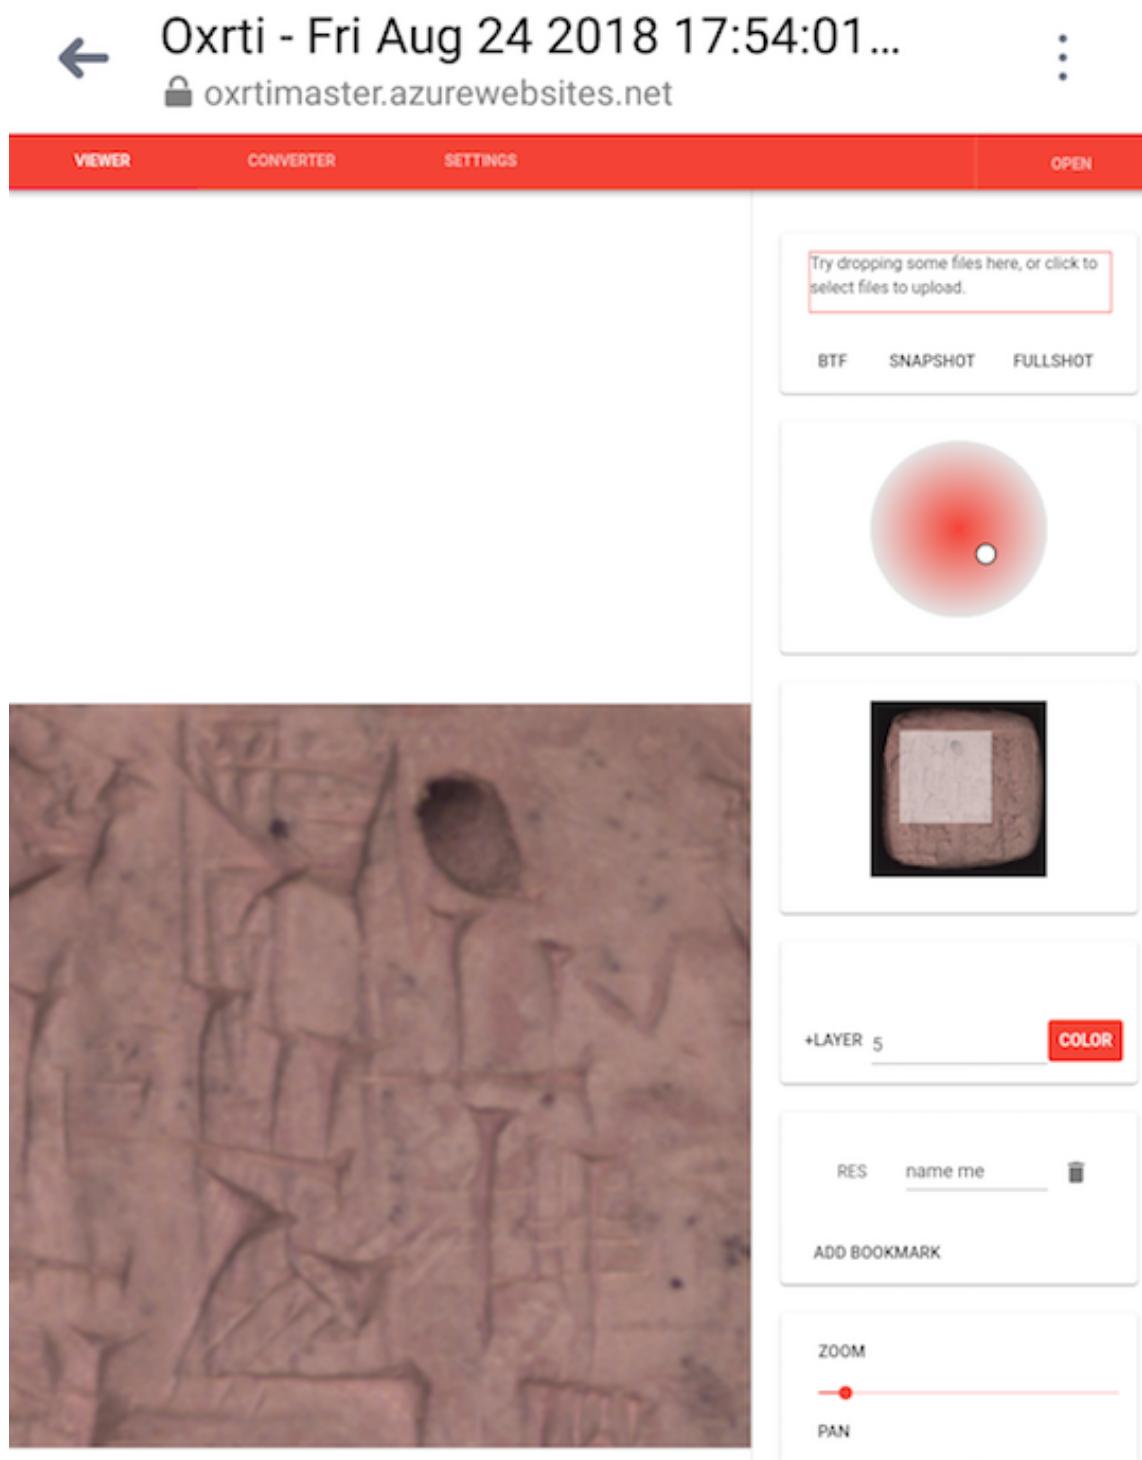


Figure 23: Implementation running on a Huawei P10 mobile phone with the Chrome browser. Opening a BTF is done through clicking the upload area.

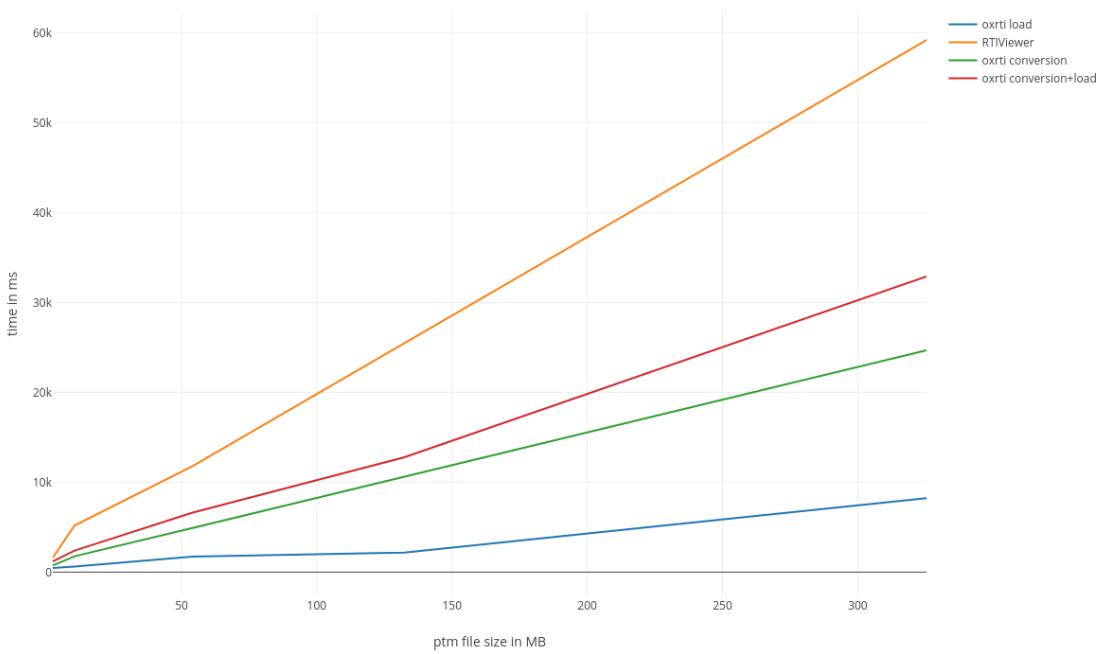


Figure 24: Orange is the time the RTIViewer takes to open the PTM file of the corresponding size. Green is the time oxrti takes to convert that PTM file into BTF format. Blue is the time oxrti takes to open the corresponding BTF file. Red sums up blue and green.

6.1 Accuracy

The next thing to verify is the accuracy; it is necessary to determine whether the implemented rendering is actually giving the 'right' image. To quantify that, the Root Mean Square Error as presented by Happa and Gogosio[21] was picked, as the RTIViewer version can be considered the gold standard for this comparison, the pixel dimensions are equal and all pixel values are calculated with no cross interactions. The formulae used is:

$$RMSE(X, \bar{X}) = \sqrt{\frac{\sum_{i=1}^N (X_i - \bar{X}_i)^2}{N}}$$

Where X are the reference pixels, \bar{X} are the exported pixels, and N represents the total number of pixels. As each image is an RGB png with 8-bit per channel; this calculation is done for each channel. Pixel values are in range [0, 255]. The images to compare are generated by setting equal light positions within the LightControl plugin and the RTIViewer and then using the respective export function. For the oxrti implementation the 'fullshot' export was used, which is exporting after base node rendering in full resolution, without any further transformation. For the RTIViewer the 'Snapshot' button was used. Due to some bug inside the the RTIViewer implementation some of the exported images have the wrong dimensions, e.g. the *Mummy_RGB* is exported as 583x999 pixels, even though it has 583x1000 pixels. These files were only compared visually. For *tablet2_LRGB* and *Warrior_RGB* the right dimensions were exported by the RTIViewer. Oxrti exports the right sizes in all cases. Having one RGB and LRGB file covers all currently implemented Visuals comparisons are shown in Figure 25 and Figure 26, the results are tabulated in Table 2. Visual inspection reveals that the images virtually look the same and the RMSE is confirming this, with less than one discrete step mean difference between the versions per channel. A comparison of a completely black with a completely white image would show a RMSE of 15.96 per channel. The small differences are likely related to shader architecture. The RTIViewer implementation can operate with pure integers, whereas the oxrti shaders use floats and their inaccuracies in addition to some potential texture filtering happening internally.

File	RMSE(R)	RMSE(G)	RMSE(B)	SUM
<i>tablet2_LRGB</i>	0.649	0.677	0.674	2.000
<i>Warrior_RGB</i>	0.627	0.612	0.610	1.849

Table 2: RMSE for Accuracy Comparison

6.2 Rollouts and Testing

All versions talked about in section 5.10 are readily available for public use. Some external testing of these releases has began with participants from the Oxford Centre for the Study of Ancient Documents (CSAD) and the Institute of Archaeology. Initial feedback indicates approval of the implementation and will be reported on in detail at a later point.



(a) Warrior from RTIViewer

(b) Warrior from oxrti

(c) Mapped diff

Figure 25: Comparison of RGB rendering between oxrti and RTIViewer. Light position set in both applications to $x = 0.7071, y := -0.7071$, so a light from bottom right. The difference is mapped from $[0, 1, 2]$ to $[0, 127, 255]$ for visualisation.



(a) Tablet from RTIViewer

(b) Tablet from oxrti

(c) Mapped diff

Figure 26: Comparison of LRGB rendering between oxrti and RTIViewer. Light position set in both applications to $x = -1, y = 0$, so a light from the left. The difference is mapped from $[0, 1, 2]$ to $[0, 127, 255]$ for visualisation.

7 Discussion

This section discusses the achieved results and proposes future avenues of research.

7.1 Novelties

Though some functionality of current RTI viewers could not be replicated during this project (mainly RTI file format support and, many novel things were achieved:

- Free Software first, the source code[1], the file format specification and accompanying documentation[3] are freely available online from first day on.
- Full plugin orientation making it the most flexible and modifiable implementation. The development speed picked up considerably feature wise the longer the project went, as most hooks were defined and functions could be stitched together to achieve fast turnarounds.
- Full web compatibility without local/native dependencies for preprocessing, enabling a speedy and comfortable workflow, with no requirements of dropping into a command line anywhere.
- Most transparent and flexible file format for RTI, while being smaller and faster to load.
- Paint plugin, making round trips with graphical editing software unnecessary.
- Flexible shader based architecture for improved performance.

7.2 Future Work

The future work can be split into two parts. Improvements of the current system, including better performance and bug fixes, and further extensions with new functionality. The following improvements are currently indicated (in order of importance):

- Support of RTI files. Requiring the addition of a new parser and some new rendering base nodes.
- Support for some rendering modes like specular enhancement and others as presented by Palma et al.[33].

For new features three areas are interesting:

Technical features:

Compression: Support for the newly (June 2018) developed techniques by Ponchio et al. to compress and enhance RTI data for web applications[35]

Animation support: The state-driven architecture would easily support some kind of tweening for light positions, zoom, etc. The image export could be extended to support video output.

3D Engines: Integration of some full blown 3D engine to support virtual scenes, with the base nodes providing underlying textures for some objects. Would require a reworking of the caching behaviour

HDR images: The BTF file format can easily support higher resolution and higher bit-depth pictures.

WebGL 2 integration with features like texture arrays could simplify some calculations. Hold up by WebGL2 compatibility of common browser though.

Rendering tree instead of rendering stack. The biggest performance problem currently is the GPU's fill rate. A more flexible tree would require fewer drawn pixels per user input, resulting in better performance.

Shader collation: Currently there is no easy exchange between different shader's code. Internal support/use of glslify[19] could help with that and potentially allow the dynamic construction of shaders at runtime, purely depending on the *manifest.json* and not on handwritten support for all formats.

Being a web application some collaborative features come to mind, if the addition of a server is considered:

Online object registry: Allowing captured objects to be registered under some hash to uniquely identify them. Potential hosting of the corresponding data file. Standard features like registration, login, privacy settings for objects.

Shared items : Bookmarks, Notes, Layers, etc. could all be exchanged online and potentially automatically synced between machines. The mobx-state-tree in the background is automatically generating patch files and all these objects have unique ids, making realtime exchange quite feasible.

Machine Learning: Having a comprehensive database of crowd sourced labelled data - e.g. labels for single pictograms of captured cuneiforms - could be the base for some machine learning on the RTI data. Also specific light configurations could be rated in their usefulness for analysing specific parts, feeding into an auto lightning button.

The success of a shader based viewer indicates that also the RTI building part could be ready for some improvements through new oxrti plugins.

Shader based: A new builder could use the GPU to parallelize the RTI construction process, which currently are single-threaded and CPU bound. A modern data set likely consist of more than 64 differently lighted images, which would not be transferable at the same time to the (WebGL controlled) GPU directly, but two options could be considered: Implementing the building part inside the Electron app and its local process, which has full access to a machine's GPU without WebGL limitations, but still controllable by the plugin architecture through inter-process-communication. Alternatively the 'stack of images'

could instead be split along the z-axis, as each texel only requires all its values for all light positions and no values of neighbouring texels, allowing for good parallelization.

Configuration based: The current RTIBuilder (also provided by CHI[11]) does an light angle estimation per captured image. With RTI domes these angles are fixed though could just be given to the software for better accuracy.

Video based: The capture and transfer of a lot of single images, which are nearly the same (as the light will only change slightly between each picture) could be accelerated by instead capturing a video. The WebGL support for video textures could be an entry point.

7.3 Community Outreach and Onboarding

The most important next step though is community outreach and onboarding as the nicest features only become useful if they are used. Outreach will happen in two different communities. The implementation core is providing a novel and modern infrastructure for any kind of plugin based software, and as such will be taken by the author to the JavaScript and TypeScript community. But more important is the RTI community. Initial outreach on social media platforms generated interest within the Classics community; the next steps are a more formal introduction of the software at the CHI forums[5], as they are frequented by many RTI users and a publication targeted at researchers in the field of archaeology and/or conservation.

8 Conclusion

The project succeeded in developing a modern, exploratory, more ‘industrialized’, web-based RTI viewer software with improved performance and new functionality; and in outlining future potential for novelties in applications around RTI. Whilst the author will ensure continued support for the software, the future plurality of developers and users of this software is of primary importance.

This is the specification of the BTF fileformat, as of version 1.0 on August 28, 2018. It was developed by co-supervisor Stefano Gogioso, with input from and extensions by the author of the enclosing thesis in relation to the oxrti viewer.

A BTF File Format

This section describes the BTF file format. The aim of this file format is to provide a generic container for BTF data to be specified using a variety of common formats. Files shall have the `.btf.zip` extension.

A.1 File Structure

A BTF file is a ZIP file containing the following:

- A **manifest** file in JSON format, named `manifest.json`. The manifest contains all information about the BRDF/BSDF model being used, including the names for the available **channels** (e.g. R, G and B for the 3-channel RGB), the names of the necessary **coefficients** (e.g. bi-quadratic coefficients) and the **image file format** for each channel.
- A single folder named **data**, with sub-folders having names in 1-to-1 correspondence with the channels specified in the manifest.
- Within each channel folder, greyscale image files having names in 1-to-1 correspondence with the coefficients specified in the manifest, each in the image file format specified in the manifest for the corresponding channel. For example, if one is working with RGB format (3-channels named R, G and B) in the PTM model (five coefficients a_2 , b_2 , a_1 , b_1 and c , specifying a bi-quadratic) using 16-bit greyscale bitmaps, the file `/data/B/a2.bmp` is the texture encoding the a_2 coefficient for the blue channel of each point in texture space.
- The datafiles are all in reversed scanline order (meaning from bottom to top), to keep aligned with the original PTM format and allow easier loading into WebGL.

In case of usage with the oxrti viewer, following files can be present in addition to those mentioned above:

- A **Layers** file in JSON format, named `oxrti_layers.json`. In addition a **layers** folder should be present, containing the textures for the layers.
- A **State** file in JSON format, named `oxrti_state.json`.

A.2 Manifest

The manifest for the BTF file format is a JSON file with root dictionary. The **root** element has two mandatory child elements: one named **data**, and one named **name** with the option of additional child elements (with different names) left open to future extensions of the format.

- The **name** element is a string with a name of the contained object.
- The **data** element has for entries, named **width**, **height**, **channels** and **channel-model**. The **width** and **height** attributes have values in the positive integers describing the dimensions of the BTDF. The **channel-model** attribute has value a non-empty alphanumeric string uniquely identifying the BRDF/BSDF colour model used by the BTF file (see the Options section below). The **channels** element has an arbitrary amount of named **channel** entries, according to the **channel-model**.
- Additionally the **data** element has one untyped entry named **formatExtra**, where format implementation specific data can be stored.
- Each **channel** has an **coefficients** child consisting of an arbitrary number of **coefficient** entries, as well as one **coefficient-model** attribute. The **coefficient-model** attribute has value a non-empty alphanumeric string uniquely identifying the BRDF/BSDF approximation model used by the BTF file (see the Options section below).
- Each **coefficient** element has one attribute: **format**. The **format** attribute has value a non-empty alphanumeric string uniquely identifying the image file format used to store the channel values (see the Options section below).

A.3 Textures

Each image file `/data/CHAN/COEFF.EXT` has the same dimensions specified by the **width** and **height** attributes of the **data** element in the manifest, and is encoded in the greyscale image file format specified by the **format** attribute of the unique **coefficient** element with attribute **name** taking the value **COEFF** (the extension `.EXT` is ignored). The colour value of a pixel (u, v) in the image is the value for coefficient **COEFF** of channel **CHAN** in the BRDF/BSDF for point (u, v) , according to the model jointly specified by the values of the attribute **model** for element **channels** (colour model) and the attribute **model** for element **coefficients** (approximation model).

In case of collated coefficients, the texture will not be greyscale, but RGB or RGBA and the coefficients are split into the RGBA channels in order of their appearance inside the **coefficient**'s name.

A.4 Options

At present, the following values are defined for attribute `channel-model` of element `channels`.

- `RGB`: the 3-channel RGB colour model, with channels named `R`, `G` and `B`. This colour model is currently under implementation.
- `LRGB`: the 4-channel LRGB colour model, with channels named `L`, `R`, `G` and `B`. This colour model is currently under implementation.
- `SPECTRAL`: the spectral radiance model, with an arbitrary non-zero number of channels named either all by wavelength (format `--nm`, with `--` an arbitrary non-zero number) or all by frequency format `--Hz`, with `--` an arbitrary non-zero number. This colour model is planned for future implementation.

At present, the following values are defined for attribute `coefficient-model` of coefficient elements, where the ending character `*` is to be replaced by an arbitrary number greater than or equal to 1.

- `flat`: flat approximation model (no dependence on light position). This approximation model is currently under implementation.
- `RTIpoly*`: order-* polynomial approximation model for RTI (single view-point BRDF). This approximation model is currently under implementation.
- `RTIharmonic*`: order-* hemispherical harmonic approximation model for RTI (single view-point BRDF). This approximation model is currently under implementation.
- `BRDFpoly*`: order-* polynomial approximation model for BRDFs. This approximation model is planned for future implementation.
- `BRDFharmonic*`: order-* hemispherical harmonic approximation model for BRDFs. This approximation model is planned for future implementation.
- `BSDFpoly*`: order-* polynomial approximation model for BSDFs. This approximation model is planned for future implementation.
- `BSDFharmonic*`: order-* spherical harmonic approximation model for BSDFs. This approximation model is planned for future implementation.

At present, the following values are defined for attribute `format` of elements tagged `coefficient`, where the ending character `*` is the bit-depth, to be replaced by an allowed positive multiple of 8.

- `BMP*`: greyscale BMP file format with the specified bit-depth (8, 16, 24 or 32). Support for this format is currently under implementation.
- `PNG*`: PNG file format encoding the specified bit-depth (8, 16, 24, 32, 48 or 64). Support for this format is currently under implementation. Different PNG colour options are used to support different bit-depths:
- `Grayscale` with 8-bit/channel to encode 8-bit bit-depth.

- `Grayscale` with 16-bit/channel to encode 16-bit bit-depth.
- `Truecolor` with 8-bit/channel to encode 24-bit bit-depth.
- `Truecolor and alpha` with 8-bit/channel to encode 32-bit bit-depth.
- `Truecolor` with 16-bit/channel to encode 48-bit bit-depth.
- `Truecolor and alpha` with 16-bit/channel to encode 64-bit bit-depth.

A.5 Layers

The layer file is a JSON file with a root array, containing a sorted list of layer information dictionaries, which consist of 3 attributes each:

- A `name` JSON string. It is also the prefix of the layer's texture inside the `layers` folder, where the corresponding `name.png` file resides.
- A `texture` object with no further attributes, to be replaced at runtime with the actual texture data.
- A `id` JSON string, to be changed each time a layer texture is exported.

A.6 State

The state file is a JSON file with oxrти implementation details and is unlikely to be of further use for other software, it basically stores the state of all loaded plugins.

References

- [1] ksjogo. *oxrti: Reflectance Transformation Imaging Toolset*. Aug. 15, 2018. URL: <https://github.com/ksjogo/oxrti> (visited on 08/17/2018).
- [2] ksjogo. *oxrti_data: public data test files for the oxrti project*. Aug. 26, 2018. URL: https://github.com/ksjogo/oxrti_data (visited on 08/26/2018).
- [3] ksjogo. *oxrti_thesis*. Aug. 28, 2018. URL: https://github.com/ksjogo/oxrti_thesis (visited on 08/28/2018).
- [4] *Alternative syntax madness · Issue #487 · mobxjs/mobx-state-tree*. GitHub. URL: <https://github.com/mobxjs/mobx-state-tree/issues/487> (visited on 08/21/2018).
- [5] CHIForums: All Viewers. Cultural Heritage Imaging Forums. URL: <http://forums.culturalheritageimaging.org/forum/28-all-viewers/> (visited on 08/28/2018).
- [6] *Classes · TypeScript*. URL: <https://www.typescriptlang.org/docs/handbook/classes.html> (visited on 08/20/2018).
- [7] *classy-mst: ES6-like syntax for mobx-state-tree*. Aug. 11, 2018. URL: <https://github.com/charto/classy-mst> (visited on 08/13/2018).
- [8] Library of Congress. *Polynomial Texture Map (PTM) File Format*. June 14, 2018. URL: <https://www.loc.gov/preservation/digital/formats//fdd/fdd000487.shtml> (visited on 08/10/2018).
- [9] Library of Congress. *Reflectance Transformation Imaging (RTI) File Format*. June 9, 2018. URL: <https://www.loc.gov/preservation/digital/formats//fdd/fdd000486.shtml#notes> (visited on 08/10/2018).
- [10] Massimiliano Corsini, Prabath Gunawardane, and Carla Schroer. “RTI Format – Draft 0.9”. In: *Cultural Heritage Imaging Forum* (2010). URL: http://forums.culturalheritageimaging.org/index.php?app=core&module=attach§ion=attach&attach_id=81.
- [11] *Cultural Heritage Imaging | Process: RTIBuilder Download*. URL: http://culturalheritageimaging.org/What_We_Offer/Downloads/Process/ (visited on 08/28/2018).
- [12] *Cultural Heritage Imaging | Reflectance Transformation Imaging (RTI)*. URL: <http://culturalheritageimaging.org/Technologies/RTI/> (visited on 08/24/2018).
- [13] *Cultural Heritage Imaging | View: RTIViewer Download*. URL: http://culturalheritageimaging.org/What_We_Offer/Downloads/View/ (visited on 08/26/2018).
- [14] *Decorators · TypeScript*. URL: <https://www.typescriptlang.org/docs/handbook/decorators.html> (visited on 08/20/2018).
- [15] *electron: Build cross-platform desktop apps with JavaScript, HTML, and CSS*. Aug. 21, 2018. URL: <https://github.com/electron/electron> (visited on 08/21/2018).
- [16] *electron-userland/electron-webpack: Scripts and configurations to compile Electron applications using webpack*. URL: <https://github.com/electron-userland/electron-webpack> (visited on 08/21/2018).

- [17] *Generics · TypeScript*. URL: <https://www.typescriptlang.org/docs/handbook/generics.html> (visited on 08/20/2018).
- [18] *gl-react cookbook*. URL: <https://gl-react-cookbook.surge.sh/hellogl> (visited on 08/21/2018).
- [19] *glslify/glslify : A node.js-style module system for GLSL!* Aug. 27, 2018. URL: <https://github.com/glslify/glslify> (visited on 08/28/2018).
- [20] *GNU Emacs - GNU Project*. URL: <https://www.gnu.org/software/emacs/> (visited on 08/17/2018).
- [21] Stefano Gogioso and Jassim Happa. “PBR: Image Synthesis Validation”. Physical Based Rendering course, Oxford University., Nov. 24, 2017.
- [22] *HP Labs : Research : Polynomial texture mapping : Downloads*. URL: <http://www.hpl.hp.com/research/ptm/downloads/download.html> (visited on 08/27/2018).
- [23] Takahiro Ethan Ikeuchi. *React Stateless Functional Component with TypeScript*. Medium. Apr. 5, 2017. URL: https://medium.com/@ethan_ikt/react-stateless-functional-component-with-typescript-ce5043466011 (visited on 08/20/2018).
- [24] *InscriptiFact*. URL: <http://www.inscriptifact.com/index.shtml> (visited on 08/27/2018).
- [25] *InscriptiFact :: Instructions ::* URL: <http://ruth.usc.edu:7060/index.jsp> (visited on 08/27/2018).
- [26] *JSZip*. URL: <https://stuk.github.io/jszip/> (visited on 08/21/2018).
- [27] Lindsay MacDonald and Stuart Robson. “Polynomial texture mapping and 3d representations”. In: *International Archives of Photogrammetry, Remote Sensing and Spatial Information Sciences 38* 5 (2010), p. 6.
- [28] Thomas Malzbender, Dan Gelb, and Hans Wolters. “Polynomial texture maps”. In: *Proceedings of the ACM SIGGRAPH Conference on Computer Graphics*. Vol. 2001. Jan. 1, 2001, pp. 519–528. DOI: [10.1145/383259.383320](https://doi.org/10.1145/383259.383320).
- [29] *material-ui: React components that implement Google’s Material Design*. Aug. 21, 2018. URL: <https://github.com/mui-org/material-ui> (visited on 08/21/2018).
- [30] *mobx: Simple, scalable state management*. Aug. 13, 2018. URL: <https://github.com/mobxjs/mobx> (visited on 08/13/2018).
- [31] *mobx-state-tree: Model Driven State Management*. Aug. 20, 2018. URL: <https://github.com/mobxjs/mobx-state-tree> (visited on 08/21/2018).
- [32] *OxRTI SVG Viewer*. URL: <https://c14.arch.ox.ac.uk/oxrti/OxRTIViewer.html?file=https://c14.arch.ox.ac.uk/oxrti/example/info.json> (visited on 08/27/2018).
- [33] Gianpaolo Palma et al. “Dynamic shading enhancement for reflectance transformation imaging”. In: *JOCCH* 3 (Sept. 1, 2010), p. 6. DOI: [10.1145/1841317.1841321](https://doi.org/10.1145/1841317.1841321).
- [34] Kathryn Piquette and Charles Crowther. *Developing a Reflectance Transformation Imaging (RTI) System for Inscription Documentation in Museum Collections and the Field: Case studies on ancient Egyptian and Classical material*. 2011. URL: <https://pdfs.semanticscholar.org/presentation/c46c/18dc51915934b9f51bb86d64d590068ee275.pdf>.

- [35] Federico Ponchio, Massimiliano Corsini, and Roberto Scopigno. “A Compact Representation of Relightable Images for the Web”. In: *Proceedings of the 23rd International ACM Conference on 3D Web Technology*. Web3D ’18. New York, NY, USA: ACM, 2018, 1:1–1:10. ISBN: 978-1-4503-5800-2. DOI: 10.1145/3208806.3208820. URL: <http://doi.acm.org/10.1145/3208806.3208820> (visited on 08/28/2018).
- [36] *React - A JavaScript library for building user interfaces*. URL: <https://reactjs.org/index.html> (visited on 08/13/2018).
- [37] *Redux DevTools*. URL: <https://chrome.google.com/webstore/detail/redux-devtools/lmhkpmbekcpmknklioeibfkpmffibljd> (visited on 08/22/2018).
- [38] *Reflectance Transformation Imaging - WebRTIViewer*. URL: <http://vcg.isti.cnr.it/rte/webviewer.php> (visited on 08/27/2018).
- [39] Gaëtan Renaudeau. *gl-react – React library to write and compose WebGL shaders*. Aug. 13, 2018. URL: <https://github.com/gre/gl-react> (visited on 08/13/2018).
- [40] Arian Stolwijk. *pngjs: Pure JavaScript PNG decoder*. July 20, 2018. URL: <https://github.com/arian/pngjs> (visited on 08/21/2018).
- [41] *The WebGL API: 2D and 3D graphics for the web*. MDN Web Docs. URL: https://developer.mozilla.org/en-US/docs/Web/API/WebGL_API (visited on 08/21/2018).
- [42] *TSLint*. URL: <https://palantir.github.io/tslint/> (visited on 08/17/2018).
- [43] *TSLint - Visual Studio Marketplace*. URL: <https://marketplace.visualstudio.com/items?itemName=eg2 tslint> (visited on 08/17/2018).
- [44] *TypeScript is a superset of JavaScript that compiles to clean JavaScript output*. Aug. 13, 2018. URL: <https://github.com/Microsoft/TypeScript> (visited on 08/13/2018).
- [45] *Visual Studio Code - Code Editing. Redefined*. URL: <http://code.visualstudio.com/> (visited on 08/17/2018).
- [46] *WebGL Specification*. URL: <https://www.khronos.org/registry/webgl/specs/1.0/> (visited on 08/21/2018).
- [47] *WebGL Stats*. URL: <http://webglstats.com/> (visited on 08/21/2018).
- [48] *WebGL Stats Texture Units*. URL: http://webglstats.com/webgl/parameter/MAX_TEXTURE_IMAGE_UNITS (visited on 08/21/2018).
- [49] *WebGL tutorial*. MDN Web Docs. URL: https://developer.mozilla.org/en-US/docs/Web/API/WebGL_API/Tutorial (visited on 08/21/2018).
- [50] *webpack/webpack: A bundler for javascript and friends. Packs many modules into a few bundled assets. Code Splitting allows to load parts for the application on demand. Through "loaders," modules can be CommonJs, AMD, ES6 modules, CSS, Images, JSON, Coffeescript, LESS, ... and your custom stuff.* URL: <https://github.com/webpack/webpack> (visited on 08/21/2018).